

MATH2269 APPLIED BAYESIAN STATISTICS - Bayesian Logistic Regression for Stroke Prediction: Analyzing Risk Factors, Gender Differences, and Health Conditions - Assignment III

Adeyinka K. Freeman (S3960988)

October 24, 2024

Table of Contents

1.0. INTRODUCTION	4
1.0.1. Background	4
1.0.2. Current Study	4
1.1. Loading all required Libraries	5
1.2. Reading CSV to Dataframe	7
1.3. Preprocessing our Dataframe	8
1.3.1 Structure of the Stroke Dataset	8
1.3.2 Data Cleaning and Imputation for BMI in the Stroke Dataset	9
1.3.3 Summary Statistics for the Preprocessed Stroke Dataset	10
1.3.4. Initial Data Exploration	11
1.4. Dependent and Independent Variables in the Bayesian Model	15
1.4.2.1 Observation & Test of Skewness - Independent Variables	16
1.4.4 Test of correlation - Independent Variables	19
1.4.5 Scatter Plot - Independent Variables	21
1.5 Model Specification and Justification - Bayesian Logistic Regression	23
1.5.1 Justification for the Model - Bayesian Logistic Regression	24
1.5.2 Encoding Categorical Variables	24
1.5.2.1. Identifying cardinality	24
1.5.2.2. One-hot encoding of Categorical variables	25
1.5.2.3. Defining Predictor Variables for our Model	26
1.5.2.3. Defining Predictor Variables and Preparing Sample Data for Bayesian Analysis	27
1.6 Mathematical Model Equation - Bayesian Logistic Regression	30
1.6.1 Bayesian Logistic Regression Model	30
1.7 Specification of the Prior Distribution	33
1.7.1 JAGS Model Diagram	33
1.7.2 JAGS Model Diagram - Using DiagrammeR in R	34
1.7.3 JAGS Model Diagram - Prior Specification in JAGS	37
1.8. Analyze Posterior Distributions	43
1.9 Robust Bayesian Logistic Regression: Sensitivity Analysis of Prior Distributions in Stroke Prediction Models	54

1.10. Balance Outcome Accuracy Between Positive and Negative Outcome Predictions	68
1.11. Conduct Predictive Checks	72
1.12. CONCLUSION AND FURTHER STUDY	75
1.12.1 CONCLUSION.....	75
1.12.2 Recommendations for Further Study	75
2. REFERENCES.....	77
Figure 1: Univariate Analysis Plot for Numeric Variables	12
Figure 2: Univariate Analysis Plot for Categorical Variables	13
Figure 3: Stroke Proportion by Categorical Variables	14
Figure 4: Distribution of Stroke (Dependent Variables).....	15
Figure 5: Histogram of Age (Predictor).....	17
Figure 6: Histogram of Average Glucose Level (Predictor)	17
Figure 7: Histogram of BMI (Predictor).....	18
Figure 8: Correlation Matrix Heatmap	20
Figure 9: Stroke vs Age.....	22
Figure 10: Stroke vs Average Glucose Level	22
Figure 11: Stroke vs BMI	23
Figure 12: Mathematical Function for our proposed Model	31
Figure 13: Diagrammatic Outlook for our Proposed Model	34
Figure 14: stroke_model_visualization	36
Figure 15: Posterior Analysis - Post Modeling.....	46
Figure 16: Density - Beta Coefficient (Predictors)	57

1.0. INTRODUCTION

Stroke is a significant public health issue and a leading cause of death globally, responsible for approximately 11% of total deaths according to the World Health Organization (2021).

Understanding the various factors contributing to stroke risk is essential for developing effective prevention strategies. This analysis utilizes a dataset that includes diverse individual characteristics, enabling a thorough examination of how these factors correlate with stroke incidence.

By employing a Bayesian modeling approach, particularly Bayesian logistic regression, this study aims to predict the likelihood of stroke based on these input parameters, thereby aiding healthcare professionals in identifying and mitigating risk factors among patients.

1.0.1. Background

Stroke is a complex medical condition influenced by a myriad of factors, including genetics, lifestyle, and pre-existing health conditions. The burden of stroke varies considerably across populations, highlighting the need for targeted research that addresses local risk factors (Feigin et al., 2014).

Statistical modeling plays a crucial role in understanding the intricate relationships between various predictors and stroke risk. The dataset analyzed in this study includes 5,110 observations, with 12 attributes capturing key demographic and health-related information.

This comprehensive dataset facilitates an in-depth exploration of how factors such as age, gender, hypertension, and lifestyle choices contribute to the risk of stroke. By employing Bayesian logistic regression, the analysis aims to identify significant predictors of stroke, ultimately contributing to better risk assessment and preventive healthcare strategies.

1.0.2. Current Study

This study aims to employ Bayesian logistic regression to explore the various factors contributing to the risk of stroke among different populations. We will specifically investigate the impact of demographic variables, health conditions, and lifestyle factors on stroke incidence. The dataset's structure allows us to analyze how these variables interact, considering potential correlations and dependencies between them.

By treating demographic factors such as gender and age as key predictors and exploring their interactions with lifestyle choices and medical history, we intend to identify significant contributors to stroke risk. Furthermore, the use of Bayesian methods will enable us to incorporate prior knowledge and assess uncertainty in our estimates, providing a more robust understanding of the predictors of stroke. The insights derived

from this analysis will be instrumental in guiding public health initiatives aimed at mitigating stroke risk in vulnerable population.

1.1. Loading all required Libraries

```
# Define required libraries
required_libraries <- c("rjags", "runjags", "ROCR", "caret", "coda",
                        "DiagrammeRsvg", "bayestestR", "ggplot2",
                        "mcmcse", "tidyr", "dplyr", "kableExtra",
                        "rsvg", "psych", "gridExtra", "moments",
                        "DiagrammeR", "magick")

# Function to install missing libraries
install_missing_packages <- function(libs) {
  missing <- libs[!libs %in% rownames(installed.packages())]
  if (length(missing) > 0) {
    install.packages(missing)
  }
}

# Install any missing libraries
install_missing_packages(required_libraries)

# Load all required libraries
lapply(required_libraries, library, character.only = TRUE)
```

OUTPUT

```
## [1] "rjags"      "coda"      "stats"     "graphics"  "grDevices" "utils"
## [7] "datasets"  "methods"   "base"
## [1] "runjags"   "rjags"     "coda"      "stats"     "graphics"
"grDevices"
## [7] "utils"     "datasets"  "methods"   "base"
## [1] "ROCR"      "runjags"   "rjags"     "coda"      "stats"
"graphics"
## [7] "grDevices" "utils"     "datasets"  "methods"   "base"
## [1] "caret"     "lattice"   "ggplot2"   "ROCR"      "runjags"   "rjags"
## [7] "coda"      "stats"     "graphics"  "grDevices" "utils"
"datasets"
## [13] "methods"   "base"
## [1] "DiagrammeRsvg" "caret"     "lattice"   "ggplot2"
## [5] "ROCR"        "runjags"   "rjags"     "coda"
## [9] "stats"       "graphics"   "grDevices" "utils"
## [13] "datasets"   "methods"   "base"
##
## [[7]]
## [1] "bayestestR" "DiagrammeRsvg" "caret"      "lattice"
## [5] "ggplot2"     "ROCR"          "runjags"    "rjags"
```

```

## [1] "mcmcse"          "bayestestR"      "DiagrammeRsvg"  "caret"
## [5] "lattice"         "ggplot2"         "ROCR"           "runjags"
## [9] "rjags"           "coda"            "stats"          "graphics"
## [13] "grDevices"       "utils"           "datasets"       "methods"
## [17] "base"
## [1] "tidyr"           "mcmcse"          "bayestestR"     "DiagrammeRsvg"
## [5] "caret"           "lattice"         "ggplot2"        "ROCR"
## [9] "runjags"         "rjags"           "coda"           "stats"
## [13] "graphics"        "grDevices"       "utils"          "datasets"
## [17] "methods"        "base"
##
## [[11]]
## [1] "dplyr"           "tidyr"           "mcmcse"          "bayestestR"
## [5] "DiagrammeRsvg"  "caret"           "lattice"         "ggplot2"
## [9] "ROCR"           "runjags"         "rjags"           "coda"
## [13] "stats"          "graphics"        "grDevices"       "utils"
## [17] "datasets"       "methods"        "base"
## [1] "kableExtra"     "dplyr"           "tidyr"           "mcmcse"
## [5] "bayestestR"     "DiagrammeRsvg"  "caret"           "lattice"
## [9] "ggplot2"        "ROCR"           "runjags"         "rjags"
## [13] "coda"           "stats"          "graphics"        "grDevices"
## [17] "utils"          "datasets"       "methods"        "base"
##
## [[13]]
## [1] "rsvg"           "kableExtra"     "dplyr"           "tidyr"
## [5] "mcmcse"         "bayestestR"     "DiagrammeRsvg"  "caret"
## [9] "lattice"        "ggplot2"        "ROCR"           "runjags"
## [13] "rjags"          "coda"           "stats"          "graphics"
## [17] "grDevices"      "utils"          "datasets"       "methods"
## [21] "base"
##
## [[14]]
## [1] "psych"          "rsvg"           "kableExtra"     "dplyr"
## [5] "tidyr"          "mcmcse"         "bayestestR"     "DiagrammeRsvg"
## [9] "caret"          "lattice"        "ggplot2"        "ROCR"
## [13] "runjags"        "rjags"          "coda"           "stats"
## [17] "graphics"       "grDevices"      "utils"          "datasets"
## [21] "methods"       "base"
##
## [[15]]
## [1] "gridExtra"      "psych"          "rsvg"           "kableExtra"
## [5] "dplyr"          "tidyr"          "mcmcse"         "bayestestR"
## [9] "DiagrammeRsvg"  "caret"          "lattice"        "ggplot2"
## [13] "ROCR"           "runjags"        "rjags"          "coda"
## [17] "stats"          "graphics"       "grDevices"      "utils"
## [21] "datasets"       "methods"        "base"
##
## [[16]]
## [1] "moments"        "gridExtra"      "psych"          "rsvg"
## [5] "kableExtra"     "dplyr"          "tidyr"          "mcmcse"

```

```
## [9] "bayestestR"      "DiagrammeRsvg" "caret"         "lattice"
## [13] "ggplot2"        "ROCR"          "runjags"       "rjags"
## [17] "coda"           "stats"         "graphics"      "grDevices"
## [21] "utils"          "datasets"      "methods"       "base"
##
## [[17]]
## [1] "DiagrammeR"      "moments"        "gridExtra"     "psych"
## [5] "rsvg"            "kableExtra"     "dplyr"         "tidyr"
## [9] "mcmcse"          "bayestestR"     "DiagrammeRsvg" "caret"
## [13] "lattice"         "ggplot2"        "ROCR"          "runjags"
## [17] "rjags"           "coda"           "stats"         "graphics"
## [21] "grDevices"       "utils"          "datasets"      "methods"
## [25] "base"
##
## [[18]]
## [1] "magick"          "DiagrammeR"     "moments"        "gridExtra"
## [5] "psych"           "rsvg"           "kableExtra"     "dplyr"
## [9] "tidyr"           "mcmcse"         "bayestestR"     "DiagrammeRsvg"
## [13] "caret"           "lattice"        "ggplot2"        "ROCR"
## [17] "runjags"         "rjags"          "coda"           "stats"
## [21] "graphics"        "grDevices"      "utils"          "datasets"
## [25] "methods"         "base"
```

1.2. Reading CSV to Dataframe

```
# Load utility functions (if applicable)
source("DBDA2E-utilities.R")
```

OUTPUT

```
##
## *****
## Kruschke, J. K. (2015). Doing Bayesian Data Analysis, Second Edition:
## A Tutorial with R, JAGS, and Stan. Academic Press / Elsevier.
## *****

# Set the file path for the dataset
file_path <- "FINALS/healthcare-dataset-stroke-data.csv"

# Read the CSV file into a dataframe
stroke_data <- read.csv(file_path)

# Display the original column names
print(colnames(stroke_data))
```

OUTPUT

```
## [1] "id"          "gender"      "age"
## [4] "hypertension" "heart_disease" "ever_married"
## [7] "work_type"    "Residence_type" "avg_glucose_level"
## [10] "bmi"         "smoking_status" "stroke"
```

1.3. Preprocessing our Dataframe

1.3.1 Structure of the Stroke Dataset

```
head(stroke_data)
```

OUTPUT

```
##      id gender age hypertension heart_disease ever_married work_type
## 1  9046  Male  67              0              1          Yes   Private
## 2  51676 Female  61              0              0          Yes Self-employed
## 3  31112  Male  80              0              1          Yes   Private
## 4  60182 Female  49              0              0          Yes   Private
## 5   1665 Female  79              1              0          Yes Self-employed
## 6  56669  Male  81              0              0          Yes   Private
##  Residence_type avg_glucose_level bmi smoking_status stroke
## 1             Urban          228.69 36.6 formerly smoked      1
## 2             Rural          202.21 N/A  never smoked      1
## 3             Rural          105.92 32.5  never smoked      1
## 4             Urban          171.23 34.4         smokes      1
## 5             Rural          174.12 24   never smoked      1
## 6             Urban          186.21 29 formerly smoked      1
```

```
str(stroke_data)
```

OUTPUT

```
## 'data.frame':    5110 obs. of  12 variables:
## $ id              : int  9046 51676 31112 60182 1665 56669 53882 10434
## 27419 60491 ...
## $ gender          : chr  "Male" "Female" "Male" "Female" ...
## $ age             : num  67 61 80 49 79 81 74 69 59 78 ...
## $ hypertension    : int  0 0 0 0 1 0 1 0 0 0 ...
## $ heart_disease    : int  1 0 1 0 0 0 1 0 0 0 ...
## $ ever_married     : chr  "Yes" "Yes" "Yes" "Yes" ...
## $ work_type        : chr  "Private" "Self-employed" "Private" "Private"
## ...
## $ Residence_type   : chr  "Urban" "Rural" "Rural" "Urban" ...
## $ avg_glucose_level: num  229 202 106 171 174 ...
## $ bmi              : chr  "36.6" "N/A" "32.5" "34.4" ...
## $ smoking_status   : chr  "formerly smoked" "never smoked" "never smoked"
## "smokes" ...
## $ stroke           : int  1 1 1 1 1 1 1 1 1 1 ...
```

The following observation can be made about the structure of our [data](#): -

- id: Integer - Unique identifier for each patient.
- gender: Character - Indicates the gender of the patient (e.g., “Male”, “Female”).
- age: Numeric - Age of the patient.
- hypertension: Integer - Indicates if the patient has hypertension (0 = No, 1 = Yes).
- heart_disease: Integer - Indicates if the patient has heart disease (0 = No, 1 = Yes).
- ever_married: Character - Indicates marital status (e.g., “Yes”, “No”).
- work_type: Character - Type of work (e.g., “Private”, “Self-employed”).
- Residence_type: Character - Type of residence (e.g., “Urban”, “Rural”).
- avg_glucose_level: Numeric - Average glucose level in the blood.
- bmi: Character - Body mass index (should be numeric but contains “N/A”).
- smoking_status: Character - Indicates smoking status (e.g., “formerly smoked”, “never smoked”).
- stroke: Integer - Indicates if the patient had a stroke (0 = No, 1 = Yes).

1.3.2 Data Cleaning and Imputation for BMI in the Stroke Dataset

```
# Handle missing values in bmi
stroke_data$bmi[stroke_data$bmi == "N/A"] <- NA # Replace "N/A" with NA
stroke_data$bmi <- as.numeric(stroke_data$bmi) # Convert to numeric
```

```
# Check for missing values after cleaning
missing_values <- colSums(is.na(stroke_data))
print(missing_values)
```

OUTPUT

```
##           id           gender           age           hypertension
##           0              0              0              0
## heart_disease ever_married work_type Residence_type
##           0              0              0              0
## avg_glucose_level bmi smoking_status stroke
##           0          201              0              0
```

```
# Convert character variables to factors
stroke_data$gender <- as.factor(stroke_data$gender)
stroke_data$ever_married <- as.factor(stroke_data$ever_married)
stroke_data$work_type <- as.factor(stroke_data$work_type)
stroke_data$Residence_type <- as.factor(stroke_data$Residence_type)
stroke_data$smoking_status <- as.factor(stroke_data$smoking_status)
```

```
#Imputation to handle missing values
mean_bmi <- mean(stroke_data$bmi, na.rm = TRUE) # Calculate mean excluding NA
stroke_data$bmi[is.na(stroke_data$bmi)] <- mean_bmi # Replace NA with mean
```

```
# Check the structure again to confirm changes
str(stroke_data)
```

OUTPUT

```
## 'data.frame':    5110 obs. of  12 variables:
## $ id              : int   9046 51676 31112 60182 1665 56669 53882 10434
27419 60491 ...
## $ gender          : Factor w/ 3 levels "Female","Male",...: 2 1 2 1 1 2 2
1 1 1 ...
## $ age             : num   67 61 80 49 79 81 74 69 59 78 ...
## $ hypertension    : int    0 0 0 0 1 0 1 0 0 0 ...
## $ heart_disease    : int    1 0 1 0 0 0 1 0 0 0 ...
## $ ever_married     : Factor w/ 2 levels "No","Yes": 2 2 2 2 2 2 2 1 2 2
...
## $ work_type        : Factor w/ 5 levels "children","Govt_job",...: 4 5 4 4
5 4 4 4 4 4 ...
## $ Residence_type   : Factor w/ 2 levels "Rural","Urban": 2 1 1 2 1 2 1 2
1 2 ...
## $ avg_glucose_level: num   229 202 106 171 174 ...
## $ bmi              : num   36.6 28.9 32.5 34.4 24 ...
## $ smoking_status   : Factor w/ 4 levels "formerly smoked",...: 1 2 2 3 2 1
2 2 4 4 ...
## $ stroke           : int    1 1 1 1 1 1 1 1 1 1 ...
```

1.3.3 Summary Statistics for the Preprocessed Stroke Dataset

```
# Check for missing values
summary(stroke_data)
```

OUTPUT

```
##      id            gender      age      hypertension
## Min.   : 67      Female:2994   Min.   : 0.08   Min.   :0.00000
## 1st Qu.:17741    Male  :2115   1st Qu.:25.00   1st Qu.:0.00000
## Median :36932    Other : 1     Median :45.00   Median :0.00000
## Mean   :36518                                Mean   :43.23   Mean   :0.09746
## 3rd Qu.:54682                                3rd Qu.:61.00   3rd Qu.:0.00000
## Max.   :72940                                Max.   :82.00   Max.   :1.00000
## heart_disease    ever_married      work_type      Residence_type
## Min.   :0.00000    No :1757      children      : 687      Rural:2514
## 1st Qu.:0.00000    Yes:3353     Govt_job      : 657      Urban:2596
## Median :0.00000                                Never_worked : 22
## Mean   :0.05401                                Private      :2925
## 3rd Qu.:0.00000                                Self-employed: 819
## Max.   :1.00000
## avg_glucose_level      bmi            smoking_status      stroke
## Min.   : 55.12      Min.   :10.30   formerly smoked: 885   Min.   :0.00000
## 1st Qu.: 77.25      1st Qu.:23.80   never smoked   :1892   1st Qu.:0.00000
## Median : 91.89      Median :28.40   smokes        : 789   Median :0.00000
## Mean   :106.15      Mean   :28.89   Unknown       :1544   Mean   :0.04873
```

##	3rd Qu. :114.09	3rd Qu. :32.80	3rd Qu. :0.00000
##	Max. :271.74	Max. :97.60	Max. :1.00000

1.3.4. Initial Data Exploration

Observations above show: -

- **id:** The IDs range from 67 to 72940, indicating a unique identifier for each observation.
- **gender:** Most of the entries are female (2,994), followed by males (2,115), with one observation categorized as “Other.”
- **age:** The age distribution ranges from 0.08 to 82 years, with a mean age of approximately 43.23 years. The first quartile (1st Qu.) is 25 years, and the third quartile (3rd Qu.) is 61 years.
- **hypertension:** A binary variable where 0 indicates no hypertension and 1 indicates the presence of hypertension. The mean value (0.09746) suggests that about 9.7% of the population has hypertension.
- **heart_disease:** Like hypertension, this is also a binary variable. The mean (0.05401) indicates that approximately 5.4% of the individuals have heart disease.
- **ever_married:** This variable shows that 1,757 individuals have never been married, while 3,353 have been married at some point.
- **work_type:** The largest category is “Private” (2,925), followed by “children” (687), and “Self-employed” (819). There are a few entries for “Govt_job” (657) and “Never_worked” (22).
- **Residence_type:** There are slightly more urban residents (2,596) compared to rural residents (2,514).
- **avg_glucose_level:** This variable ranges from 55.12 to 271.74, with a mean glucose level of 106.15.
- **bmi:** The Body Mass Index (BMI) ranges from 10.30 to 97.60, with a mean of approximately 28.89, indicating a mix of underweight, normal, overweight, and obese individuals.
- **smoking_status:** The categories include “formerly smoked” (885), “never smoked” (1,892), “smokes” (789), and a significant number with “Unknown” (1,544).
- **stroke:** This binary variable indicates stroke incidence, with a mean of 0.04873, suggesting that about 4.87% of the dataset experienced a stroke.

1.3.4.1. Univariate Analysis Plots - Numeric Variables

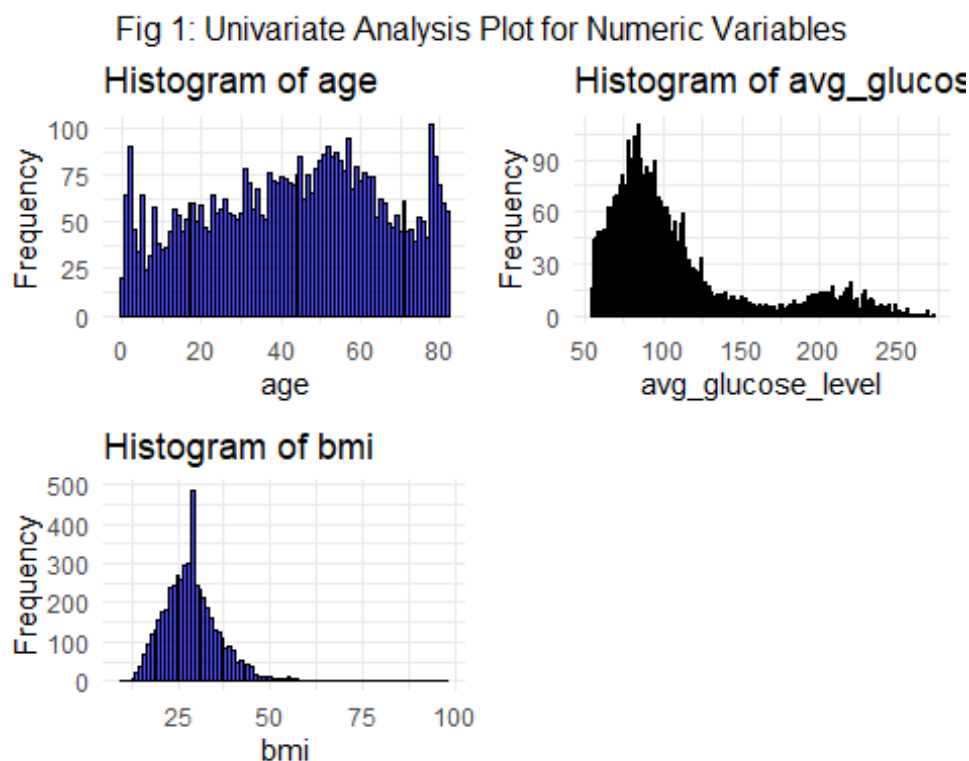
```
# Define numeric variables
num_vars <- c("age", "avg_glucose_level", "bmi")

# Generate histogram plots
fig1 <- lapply(num_vars, function(var) {
  ggplot(stroke_data, aes(x = !!sym(var))) +
    geom_histogram(binwidth = 1, fill = "blue", color = "black", alpha = 0.7)
+
  labs(title = paste("Histogram of", var), x = var, y = "Frequency") +
  theme_minimal()
})

# Arrange plots in one grid
grid.arrange(grobs = fig1, ncol = 2, top = "Fig 1: Univariate Analysis Plot
for Numeric Variables")
```

OUTPUT

Figure 1: Univariate Analysis Plot for Numeric Variables




1.3.4.2. Univariate Analysis Plot for Categorical Variables

```
# Univariate Analysis Plot for Categorical Variables
cat_vars <- c("gender", "ever_married", "smoking_status", "Residence_type",
"work_type")
```

```
# Generate bar plots
fig2 <- lapply(cat_vars, function(var) {
  ggplot(stroke_data, aes_string(x = var)) +
    geom_bar(fill = "orange", color = "black", alpha = 0.7) +
    labs(title = paste("Bar Plot of", var), x = var, y = "Count") +
    theme_minimal()
})

# Arrange plots in one grid
grid.arrange(grobs = fig2, ncol = 2, top = "Fig 2: Univariate Analysis Plot
for Categorical Variables")
```

Figure 2: Univariate Analysis Plot for Categorical Variables



1.3.4.3. Multivariate Analysis Plot for Categorical Variables

```
# Function to create bar plots with optimizations for reduced clutter
plot_categorical_vars <- function(data, variables) {
  plots <- lapply(variables, function(var) {
    ggplot(data, aes_string(x = var, fill = "as.factor(stroke)")) +
      geom_bar(position = "fill", width = 0.7) + # Adjust bar width for
spacing
    scale_y_continuous(labels = scales::percent) +
    labs(x = var, y = "Proportion", title = paste("Proportion of Stroke
by", var)) +
    scale_fill_manual(values = c("#E41A1C", "#377EB8"), labels = c("No
Stroke", "Stroke")) +
    theme_minimal(base_size = 10) + # Reduce base font size
  })
}
```

```

    theme(legend.title = element_blank(),
          plot.margin = margin(5, 5, 5, 5), # Add plot margins
          axis.text.x = element_text(angle = 45, hjust = 1)) # Rotate x-
axis labels
  })
  return(plots)
}

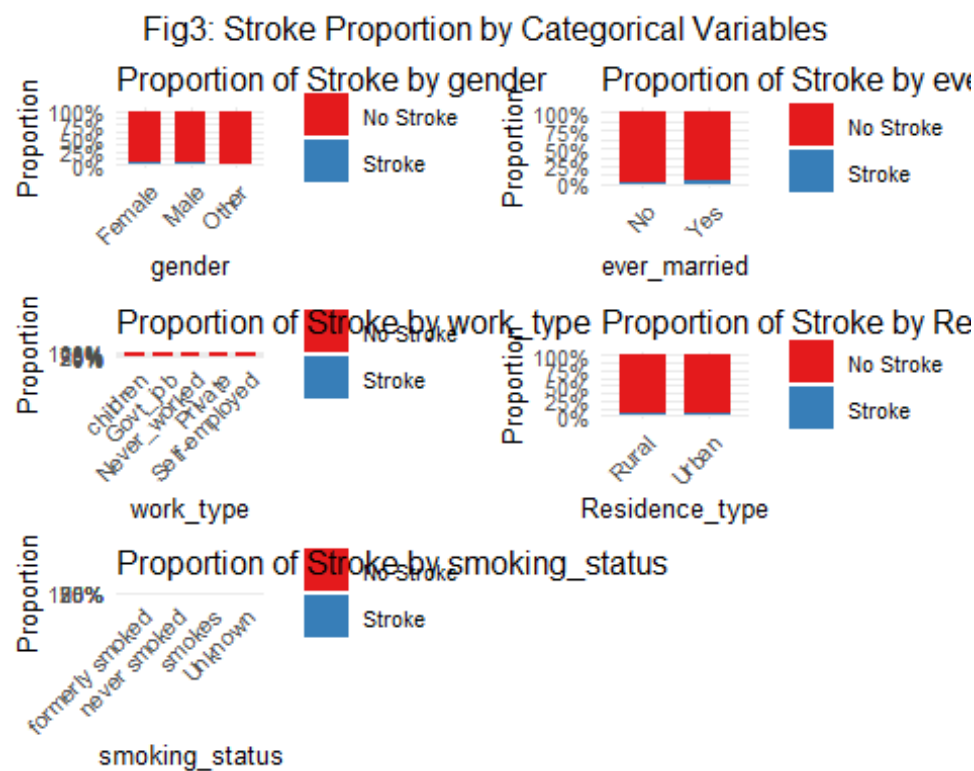
# Define categorical variables and create plots
categorical_vars <- c("gender", "ever_married", "work_type",
"Residence_type", "smoking_status")
plots <- plot_categorical_vars(stroke_data, categorical_vars)

# Arrange and display in a grid layout with increased spacing
grid.arrange(grobs = plots, ncol = 2, top = "Fig3: Stroke Proportion by
Categorical Variables")

```

OUTPUT

Figure 3: Stroke Proportion by Categorical Variables



1.4. Dependent and Independent Variables in the Bayesian Model

In our analysis, we identify and define our dependent variable and independent variables as follows:

1.4.1 Dependent Variable

- stroke: This binary variable indicates whether a patient has had a stroke (1) or not (0). This will be the outcome variable we will try to predict or explain.

1.4.1.1 Observation of our Dependent Variables

Check distribution of the dependent variable

```
table(stroke_data$stroke)
```

```
##
```

```
##    0    1
```

```
## 4861  249
```

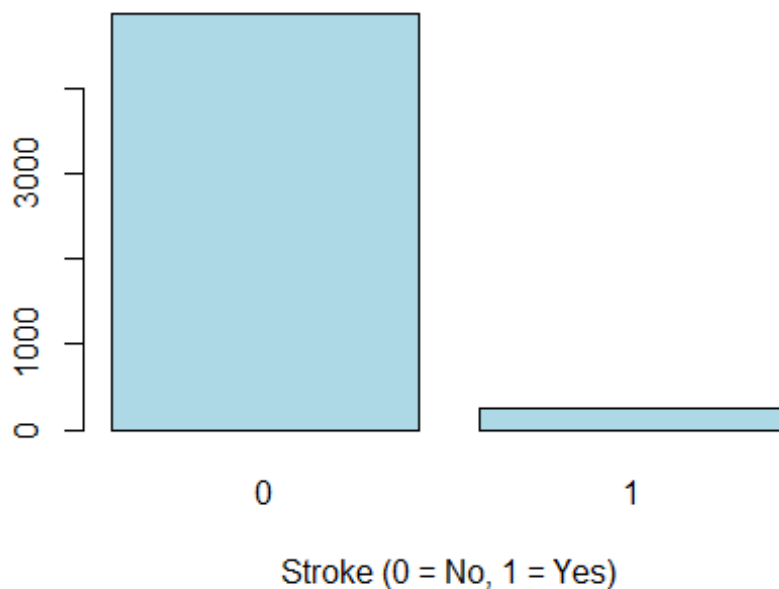
Plot a bar chart for stroke distribution

```
barplot(table(stroke_data$stroke),  
        main="Fig.4. Distribution of Stroke (Dependent Variable)",  
        xlab="Stroke (0 = No, 1 = Yes)",  
        col="lightblue")
```

OUTPUT

Figure 4: Distribution of Stroke (Dependent Variables)

Fig.4. Distribution of Stroke (Dependent Variable)



From our observation above, we can see that we have 4861 cases of no Stroke versus 249 cases of stroke occurrence.

1.4.2 Independent Variables - Bayesian Model

- gender: Categorical variable representing the gender of the individual (Male, Female, Other).
- age: Continuous variable representing the age of the individual.
- hypertension: Binary variable indicating the presence (1) or absence (0) of hypertension.
- heart_disease: Binary variable indicating the presence (1) or absence (0) of heart disease.
- ever_married: Categorical variable indicating marital status (Yes, No).
- work_type: Categorical variable representing the type of work (e.g., Private, Self-employed).
- Residence_type: Categorical variable indicating residence type (Urban, Rural).
- avg_glucose_level: Continuous variable representing average glucose level.
- bmi: Continuous variable representing body mass index.
- smoking_status: Categorical variable indicating smoking status (e.g., never smoked, formerly smoked, smokes).

1.4.2.1 Observation & Test of Skewness - Independent Variables

```
# List of numeric independent variables
independent_vars <- c("age", "avg_glucose_level", "bmi")

# Calculate skewness for independent variables
for (var in independent_vars) {
  skewness_value <- skewness(stroke_data[[var]], na.rm = TRUE)
  cat(paste("Skewness of", var, ":", skewness_value, "\n"))
}


## Skewness of age : -0.137019086639602
## Skewness of avg_glucose_level : 1.5718222973972
## Skewness of bmi : 1.0763999841401

# Set layout for individual plots
par(mfrow=c(1,1)) # Reset layout to 1 plot at a time

# Histogram with density plot for Age (Fig 2)
hist(stroke_data$age, main="Fig 5: Histogram of Age", xlab="Age",
col="lightblue", freq=FALSE)
lines(density(stroke_data$age, na.rm=TRUE), col="red", lwd=2) # Add density
plot
```


OUTPUT


Figure 5: Histogram of Age (Predictor)



```
# Histogram with density plot for Avg Glucose Level (Fig 3)
hist(stroke_data$avg_glucose_level, main="Fig 6: Histogram of Avg Glucose
Level", xlab="Avg Glucose Level", col="lightblue", freq=FALSE)
lines(density(stroke_data$avg_glucose_level, na.rm=TRUE), col="red", lwd=2)
# Add density plot
```

OUTPUT


Figure 6: Histogram of Average Glucose Level (Predictor)



```
# Histogram with density plot for BMI (Fig 4)
hist(stroke_data$bmi, main="Fig 7: Histogram of BMI", xlab="BMI",
col="lightblue", freq=FALSE)
lines(density(stroke_data$bmi, na.rm=TRUE), col="red", lwd=2) # Add density
plot
```

OUTPUT

Figure 7: Histogram of BMI (Predictor)



- Skewness of age: -0.137 A skewness value close to 0 (in this case, slightly negative) suggests that the distribution of age is approximately symmetric. The slight negative skew means that the left tail (younger individuals) is a bit longer or there are slightly more young people than older ones, but the distribution is nearly normal.
- Skewness of avg_glucose_level: 1.572 A positive skewness greater than 1 indicates a right-skewed distribution. This means that most patients have lower glucose levels, but there are a few patients with much higher glucose levels, causing the right tail to be longer. There may be outliers or extreme values in the avg_glucose_level variable that pull the distribution to the right.
- Skewness of bmi: 1.076 A skewness value of 1.076 also indicates a right-skewed distribution, although not as extreme as avg_glucose_level. This suggests that most individuals have lower or moderate BMI values, but there are some individuals with higher BMI, creating a long tail on the right side.

1.4.4 Test of correlation - Independent Variables

```
corr_matrix <- cor(stroke_data[, supply(stroke_data, is.numeric)], use =  
"complete.obs")  
print(corr_matrix)
```

OUTPUT

```
##              id          age hypertension heart_disease  
## id          1.000000000 0.003538065  0.003549615 -0.001295941  
## age          0.003538065 1.000000000  0.276397628  0.263795916  
## hypertension 0.003549615 0.276397628  1.000000000  0.108306076  
## heart_disease -0.001295941 0.263795916  0.108306076  1.000000000  
## avg_glucose_level 0.001092355 0.238171114  0.174473811  0.161857332  
## bmi          0.002998925 0.325942473  0.160188844  0.038898715  
## stroke       0.006388170 0.245257346  0.127903823  0.134913997  
##              avg_glucose_level      bmi      stroke  
## id          0.001092355 0.002998925 0.00638817  
## age          0.238171114 0.325942473 0.24525735  
## hypertension 0.174473811 0.160188844 0.12790382  
## heart_disease 0.161857332 0.038898715 0.13491400  
## avg_glucose_level 1.000000000 0.168751351 0.13194544  
## bmi          0.168751351 1.000000000 0.03894660  
## stroke       0.131945441 0.038946597 1.00000000
```

```
# Create a correlation matrix
```

```
corr_matrix <- cor(stroke_data[, supply(stroke_data, is.numeric)], use =  
"complete.obs")
```

```
# Convert the correlation matrix to a data frame
```


```
corr_df <- as.data.frame(as.table(corr_matrix))
```

```
# Create the heatmap
```

```
ggplot(corr_df, aes(Var1, Var2, fill = Freq)) +  
  geom_tile() +  
  scale_fill_gradient2(low = "blue", high = "red", mid = "white",  
                        midpoint = 0, limit = c(-1, 1),  
                        name = "Correlation") +  
  theme_minimal() +  
  theme(axis.text.x = element_text(angle = 45, hjust = 1)) +  
  labs(title = "Fig8. Correlation Matrix Heatmap", x = "Variables", y =  
"Variables")
```

OUTPUT

Figure 8: Correlation Matrix Heatmap



The correlation matrix provided displays the pairwise correlations between numeric variables in the stroke_data dataset. Here's an interpretation of the key points:

- **Id Variable:** The id variable shows a correlation of near zero with all other variables, indicating it does not influence or is not influenced by any other variable. This is expected, as id is likely just a unique identifier for the entries.
- **Age:** Age has a moderate positive correlation with both hypertension (0.276) and heart_disease (0.264), suggesting that older individuals are more likely to have hypertension and heart disease. It also shows a strong correlation with bmi (0.326), indicating a potential trend where older individuals may have higher BMI.
- **hypertension and heart_disease:** Both hypertension and heart_disease are positively correlated, indicating that individuals with hypertension are somewhat more likely to have heart disease.
- **avg_glucose_level:** The avg_glucose_level is positively correlated with hypertension (0.174) and heart_disease (0.162), suggesting a potential link between glucose levels and these conditions. A strong correlation with the log_avg_glucose_level (0.980) confirms that transforming this variable effectively maintains a strong relationship with its original form.

- BMI: BMI has a moderate correlation with age (0.326), indicating that older individuals tend to have higher BMI. The correlation with stroke is relatively weak (0.039), suggesting that while BMI may have some influence, it is not a strong predictor in this dataset.
- Log Transformed Variables: The log-transformed variables (log_avg_glucose_level and log_bmi) show strong correlations with their original counterparts (avg_glucose_level and bmi)

Observation: The variables age, hypertension, heart_disease, and avg_glucose_level have moderate positive correlations with stroke, indicating that they are potentially useful predictors. There are no variables with extremely high correlations (e.g., > 0.8), which means multicollinearity is not a concern here. Overall, the correlations suggest that age and health-related variables (hypertension, heart disease, glucose level) are more relevant predictors for stroke than BMI or id.

1.4.5 Scatter Plot - Independent Variables

```
# Define a function to create scatter plots with figure numbers
create_scatter_plot <- function(x, y, data, fig_num, x_label, y_label) {
  ggplot(data, aes_string(x = x, y = y)) +
    geom_point(alpha = 0.5, color = "blue") + # Use points with transparency
    geom_smooth(method = "lm", color = "red", se = FALSE) + # Add a Linear
    regression line
  labs(title = paste("Fig", fig_num, ":", y_label, "vs", x_label), #
    Figure number and title
    x = x_label,
    y = y_label) +
  theme_minimal() + # Use a minimal theme
  theme(plot.title = element_text(hjust = 0.5)) # Center the title
}

# Create scatter plots for different pairs with updated figure numbers
# 8. Age vs Stroke
plot8 <- create_scatter_plot("age", "stroke", stroke_data, 9, "Age",
"Stroke")


# 9. Average Glucose Level vs Stroke
plot9 <- create_scatter_plot("avg_glucose_level", "stroke", stroke_data, 10,
"Average Glucose Level", "Stroke")

# 10. BMI vs Stroke
plot10 <- create_scatter_plot("bmi", "stroke", stroke_data, 11, "BMI",
"Stroke")

# Print the plots
print(plot8)
```


OUTPUT

Figure 9: Stroke vs Age



```
print(plot9)
```


Figure 10: Stroke vs Average Glucose Level



```
print(plot10)
```

OUTPUT

Figure 11: Stroke vs BMI



Our Scatter plot above (Fig 8 - 12) validates our observations from the correlation analysis.

- Age vs. Stroke Trend: An upward trend suggests that older individuals might have a higher probability of having a stroke.
- Average Glucose Level vs. Stroke Trend: A positive trend suggests higher glucose levels are associated with a higher likelihood of stroke.
- BMI vs. Stroke Trend: A visible trend suggest a relationship. We can assume that higher BMI might correlate with increased stroke risk.

1.5 Model Specification and Justification - Bayesian Logistic Regression

Based on our results above, we would be using the Bayesian Logistic Regression the following reasons are our justifications

1.5.1 Justification for the Model - Bayesian Logistic Regression

- **Relevance:** The predictors chosen for the model (log-transformed glucose level, BMI, age, hypertension, heart disease, and lifestyle factors) are well-supported in the literature as risk factors for stroke, making this model relevant for clinical and public health research.
- **Interpretability:** Logistic regression provides easily interpretable coefficients that represent the change in the log odds of the outcome for a one-unit increase in the predictor, making it a practical choice for stakeholders in health-related fields.
- **Handling of Uncertainty:** Bayesian methods allow for the incorporation of prior knowledge and the quantification of uncertainty in parameter estimates, making them particularly suitable for medical applications where uncertainty is a crucial consideration.
- **Flexibility:** The Bayesian framework allows for the incorporation of complex hierarchical structures if needed and can adapt to various levels of data aggregation (e.g., individual-level vs. group-level data).
- **Predictive Capability:** By including both continuous and categorical variables, the model can capture non-linear relationships and interactions, which can enhance predictive performance compared to simpler models.

1.5.2 Encoding Categorical Variables

Since we intend to also include the categorical variables in our Bayesian model, we will encode it into a format that can be used. The encoding format we intend to use is the One-Hot encoding format which is a method that creates a binary column of each category in our variable. For example, for the variable `smoking_status`, we create separate columns like `smoking_status_smokes`, `smoking_status_formerly_smoked`, `smoking_status_never_smoked`, and `smoking_status_unknown`. Each of these columns would have a value of 1 if the observation belongs to that category and 0 otherwise. This action would help us also any possibilities of multicollinearity.

1.5.2.1. Identifying cardinality

```
# List of categorical variables
categorical_vars <- c("gender", "ever_married", "work_type",
"Residence_type", "smoking_status")

# Check the number of unique values for each categorical variable
cardinality <- sapply(stroke_data[categorical_vars], function(x)
length(unique(x)))

# Print the cardinality of each categorical variable
print(cardinality)
```


OUTPUT

```
##          gender  ever_married      work_type Residence_type smoking_status
##              3              2              5              2              4
```

Low Cardinality: - Ever Married and Residence Type have low cardinality (2 unique values each). These variables are suitable for binary encoding.

Moderate Cardinality: - Gender (3 unique values) and Smoking Status (4 unique values) are moderately cardinal and can be effectively one-hot encoded without significant risk of multicollinearity.

Moderate to High Cardinality: - Work Type has 5 unique values, which is still manageable for one-hot encoding.

1.5.2.2. One-hot encoding of Categorical variables

```
# Create dummy variables using model.matrix
dummy_vars <- model.matrix(~ gender + ever_married + smoking_status +
Residence_type + work_type, data = stroke_data)

# Combine dummy variables with the original dataset
stroke_data <- cbind(stroke_data, dummy_vars)

# remove the original categorical variables since no longer needed
stroke_data <- stroke_data[, !(names(stroke_data) %in% c("gender",
"ever_married", "smoking_status", "Residence_type", "work_type"))]

# Remove spaces and special characters from column names
colnames(stroke_data) <- gsub(" ", "_", colnames(stroke_data))      # Replace
spaces with underscores
colnames(stroke_data) <- gsub("-", "_", colnames(stroke_data))      # Replace
hyphens with underscores

# Optionally, ensure valid column names in case of other special characters
colnames(stroke_data) <- make.names(colnames(stroke_data), unique = TRUE)

# Print the updated column names
colnames(stroke_data)
```

OUTPUT

```
## [1] "id"          "age"
## [3] "hypertension" "heart_disease"
## [5] "avg_glucose_level" "bmi"
## [7] "stroke"        "X.Intercept."
## [9] "genderMale"    "genderOther"
## [11] "ever_marriedYes" "smoking_statusnever_smoked"
## [13] "smoking_statussmokes" "smoking_statusUnknown"
## [15] "Residence_typeUrban" "work_typeGovt_job"
```

```
## [17] "work_typeNever_worked"      "work_typePrivate"
## [19] "work_typeSelf_employed"
```

1.5.2.3. Defining Predictor Variables for our Model

```
#Set seed and sample from the stroke_data dataset
set.seed(123) # For reproducibility
sample_size <- 500

# Check if sample_size is less than or equal to the number of rows in
stroke_data
if (sample_size > nrow(stroke_data)) {
  stop("Sample size exceeds the number of available observations.")
}

# Sample the data
stroke_sample <- stroke_data[sample(nrow(stroke_data), sample_size), ]

str(stroke_sample)
```

OUTPUT

```
## 'data.frame':    500 obs. of  19 variables:
## $ id                : int  27626 29514 35602 29933 34026 55567
71917 30352 46767 22607 ...
## $ age               : num  60 43 52 5 60 76 12 57 8 41 ...
## $ hypertension      : int  0 0 0 0 0 0 0 0 0 0 ...
## $ heart_disease     : int  0 0 0 0 0 1 0 0 0 0 ...
## $ avg_glucose_level : num  266.6 97.5 107.8 86.1 207.8 ...
## $ bmi               : num  25.5 28.3 22 19 38.9 28.1 25.3 29.8 24
28.6 ...
## $ stroke            : int  0 0 0 0 0 0 0 0 0 0 ...
## $ X.Intercept.      : num  1 1 1 1 1 1 1 1 1 1 ...
## $ genderMale        : num  0 0 0 0 0 0 1 1 0 0 ...
## $ genderOther       : num  0 0 0 0 0 0 0 0 0 0 ...
## $ ever_marriedYes   : num  0 1 1 0 1 1 0 1 0 1 ...
## $ smoking_statusnever_smoked: num  1 0 0 0 1 1 1 0 0 1 ...
## $ smoking_statussmokes : num  0 0 0 0 0 0 0 0 0 0 ...
## $ smoking_statusUnknown : num  0 0 0 1 0 0 0 1 1 0 ...
## $ Residence_typeUrban : num  0 0 0 0 0 0 0 0 0 1 ...
## $ work_typeGovt_job  : num  1 0 1 0 0 0 0 0 0 0 ...
## $ work_typeNever_worked : num  0 0 0 0 0 0 0 0 0 0 ...
## $ work_typePrivate   : num  0 1 0 0 1 1 0 1 0 1 ...
## $ work_typeSelf_employed : num  0 0 0 0 0 0 0 0 0 0 ...
```

```
head(stroke_sample)
```

OUTPUT

```
##          id age hypertension heart_disease avg_glucose_level  bmi  stroke
## 2463 27626 60           0           0          266.59 25.5      0
## 2511 29514 43           0           0           97.55 28.3      0
## 2227 35602 52           0           0          107.84 22.0      0
## 526 29933 5           0           0           86.11 19.0      0
## 4291 34026 60           0           0          207.84 38.9      0
## 2986 55567 76           0           1           86.09 28.1      0
##      X.Intercept. genderMale genderOther ever_marriedYes
## 2463              1           0           0              0
## 2511              1           0           0              1
## 2227              1           0           0              1
## 526               1           0           0              0
## 4291              1           0           0              1
## 2986              1           0           0              1
##      smoking_statusnever_smoked smoking_statussmokes smoking_statusUnknown
## 2463                          1              0              0
## 2511                          0              0              0
## 2227                          0              0              0
## 526                           0              0              1
## 4291                          1              0              0
## 2986                          1              0              0
##      Residence_typeUrban work_typeGovt_job work_typeNever_worked
## 2463                   0              1              0
## 2511                   0              0              0
## 2227                   0              1              0
## 526                    0              0              0
## 4291                   0              0              0
## 2986                   0              0              0
##      work_typePrivate work_typeSelf_employed
## 2463                 0              0
## 2511                 1              0
## 2227                 0              0
## 526                  0              0
## 4291                 1              0
## 2986                 1              0
```

1.5.2.3. Defining Predictor Variables and Preparing Sample Data for Bayesian Analysis

```
# Define predictor variables
predictor_vars <- c("age", "avg_glucose_level", "bmi",
                    "genderMale", "genderOther",
                    "ever_marriedYes",
                    "smoking_statusnever_smoked",
                    "smoking_statussmokes",
                    "smoking_statusUnknown",
```

```

        "Residence_typeUrban",
        "work_typeGovt_job",
        "work_typeNever_worked",
        "work_typePrivate",
        "work_typeSelf_employed")

# Check for missing variables in data
missing_vars <- setdiff(predictor_vars, colnames(stroke_data))
if (length(missing_vars) > 0) {
  stop(paste("Missing predictor variables:", paste(missing_vars, collapse =
", ")))
}

# Check for missing variables in data
missing_vars <- setdiff(predictor_vars, colnames(stroke_data))
if (length(missing_vars) > 0) {
  stop(paste("Missing predictor variables:", paste(missing_vars, collapse =
", ")))
}

# Create a new data frame with only the predictors and the response variable
response_var <- "stroke" # Adjust if your response variable is named
differently
stroke_sample <- stroke_sample[, c(predictor_vars, response_var)]

# Check for missing values
if (anyNA(stroke_sample)) {
  warning("There are missing values in the sampled data.")
}

# Preview the sampled data
head(stroke_sample)

```

OUTPUT

```

##      age avg_glucose_level  bmi genderMale genderOther ever_marriedYes
## 2463  60          266.59 25.5          0          0          0
## 2511  43          97.55 28.3          0          0          1
## 2227  52         107.84 22.0          0          0          1
## 526   5          86.11 19.0          0          0          0
## 4291  60         207.84 38.9          0          0          1
## 2986  76          86.09 28.1          0          0          1
##      smoking_statusnever_smoked smoking_statussmokes smoking_statusUnknown
## 2463                        1                0                0
## 2511                        0                0                0
## 2227                        0                0                0
## 526                         0                0                1
## 4291                        1                0                0

```

```

## 2986          1          0          0
##      Residence_typeUrban work_typeGovt_job work_typeNever_worked
## 2463          0          1          0
## 2511          0          0          0
## 2227          0          1          0
## 526          0          0          0
## 4291          0          0          0
## 2986          0          0          0
##      work_typePrivate work_typeSelf_employed stroke
## 2463          0          0          0
## 2511          1          0          0
## 2227          0          0          0
## 526          0          0          0
## 4291          1          0          0
## 2986          1          0          0

# Output the predictor variables
predictor_vars

## [1] "age"          "avg_glucose_level"
## [3] "bmi"          "genderMale"
## [5] "genderOther"  "ever_marriedYes"
## [7] "smoking_statusnever_smoked" "smoking_statussmokes"
## [9] "smoking_statusUnknown" "Residence_typeUrban"
## [11] "work_typeGovt_job" "work_typeNever_worked"
## [13] "work_typePrivate" "work_typeSelf_employed"

```

Our predictor variable therefore shows: -

- age: Continuous variable representing the age of the individual. Older age is often associated with a higher risk of stroke.
- avg_glucose_level: Continuous variable indicating average blood glucose levels. Higher glucose levels may indicate diabetes, which is a risk factor for stroke.
- bmi: Body Mass Index, a continuous measure of body fat based on height and weight. Higher BMI may be linked to increased risk of stroke.
- gender: Categorical variables (binary) indicating gender (Male, Other). Gender can influence stroke risk, with males typically at a higher risk at younger ages.
- ever_married: Binary variable indicating marital status. Studies suggest that married individuals may have lower stroke risk due to social support.
- smoking_status: Categorical variable (binary) reflecting smoking habits. Smoking is a known risk factor for stroke.
- Residence_type: Indicates whether the individual lives in an urban or rural area, potentially impacting access to healthcare and lifestyle factors.

- **work_type:** Categorical variables indicating employment type, which may relate to stress levels and lifestyle factors associated with stroke risk.

1.6 Mathematical Model Equation - Bayesian Logistic Regression

To model the probability of experiencing a stroke (Y) based on a set of predictors (both continuous and categorical/binary variables) using Bayesian Logistic Regression.

1.6.1 Bayesian Logistic Regression Model

To model the probability of experiencing a stroke (Y) given a set of predictors, which include both continuous and categorical/binary variables. The probability can be expressed mathematically as:

1.6.1.1 Model Definition

The binary outcome variable $Y[i]$ indicates whether a stroke occurred (1) or not (0):

$$Y[i] \sim \text{Bernoulli}(\mu[i])$$

Where: - $Y[i]$: binary outcome variable indicating stroke occurrence. - $\mu[i]$: predicted probability of a stroke for observation i .

1.6.1.2 Probability Modeling

$$\mu = \text{logistic}(\beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_k x_k)$$

The predicted probability $\mu[i]$ is expressed as:

$$\mu[i] = \alpha \cdot 0.5 + (1 - \alpha) \cdot \text{logistic}(X_i \beta_k)$$

Where: - α : guessing coefficient, allowing randomness in predictions. - X_i : vector of independent variables (predictors) for observation i . - β : vector of coefficients for the predictors in X . - $\text{logistic}(z) = \frac{1}{1+e^{-z}}$: logistic function applied to $X_i \beta$.

1.6.1.3 Model Priors

The prior distributions for the parameters are specified as: - For the guessing coefficient α :

$$\alpha \sim \text{Beta}(1,9)$$

- For the intercept β_0 :

$$\beta_0 \sim \text{Normal}\left(0, \frac{1}{10}\right)$$

- For each coefficient β_k (where $j = 1, 2, \dots, N_k$):

$$\beta_k \sim \text{Normal}\left(0, \frac{1}{10}\right)$$

Where: N_k : total number of predictors in the model.

1.6.1.4 Summary of Predictors

The matrix X consists of the following predictors:

1.6.1.5 Continuous Variables

The following are the continuous variables below: -

- Age: age
- Average Glucose Level: avg_glucose_level
- BMI: bmi

Categorical Variables (encoded as dummy variables): - **Gender:** - genderMale - genderOther

- **Ever Married:** - ever_marriedYes
- **Smoking Status:** - smoking_statusnever_smoked - smoking_statussmokes - smoking_statusUnknown
- **Residence Type:** - Residence_typeUrban
- **Work Type:** - work_typeGovt_job - work_typeNever_worked - work_typePrivate - work_typeSelf_employed

Given the predictors in predictor_vars, we can express ($\mu[i]$) in a similar format to the example provided. Here's the formula for ($\mu[i]$), we apply the logistic function to a linear combination of predictors with corresponding coefficients.

Figure 12: Mathematical Function for our proposed Model

$$\mu[i] = \alpha \cdot \frac{1}{2} + (1 - \alpha) \cdot \text{logistic} \left(\beta_0 + \beta_1 \cdot \text{age} + \beta_2 \cdot \text{avg_glucose_level} + \beta_3 \cdot \text{bmi} + \beta_4 \cdot \text{genderMale} + \beta_5 \cdot \text{genderOther} + \beta_6 \cdot \text{ever_marriedYes} \right. \\ \left. + \beta_7 \cdot \text{smoking_statusnever_smoked} + \beta_8 \cdot \text{smoking_statussmokes} + \beta_9 \cdot \text{smoking_statusUnknown} + \beta_{10} \cdot \text{Residence_typeUrban} \right. \\ \left. + \beta_{11} \cdot \text{work_typeGovt_job} + \beta_{12} \cdot \text{work_typeNever_worked} + \beta_{13} \cdot \text{work_typePrivate} + \beta_{14} \cdot \text{work_typeSelf_employed} \right)$$

Where - $\left(\alpha \cdot \frac{1}{2} \right)$: The “guessing” component adds a baseline probability. - Logistic component: $\left(\text{logistic}(z) = \frac{1}{1+e^{-z}} \right)$, where $(z = \beta_0 + \sum_{k=1}^{14} \beta_k \cdot x_k)$. - (β_0) : Intercept term. - Coefficients: $(\beta_1, \beta_2, \dots, \beta_{14})$: Corresponding to each predictor in predictor_vars.

1.6.1.4 Approach and Justification of the Bayesian Logistic Regression Model Used

1.6.1.4.1. Model Structure

The model aims to predict the probability of stroke occurrence ($Y[i]$) based on various predictor variables, incorporating both continuous and categorical factors. The use of a

Bernoulli distribution for the outcome variable captures the binary nature of stroke incidence.

1.6.1.4.2. Probability Modeling

The predicted probability ($\mu[i]$) is represented through a combination of:

- A baseline guessing coefficient (α) that incorporates randomness, enabling the model to adjust predictions towards a baseline probability of 0.5 when other predictors are not strong indicators.
- The logistic function applied to a linear combination of predictors, allowing for non-linear relationships between predictors and the probability of stroke occurrence.

The model's structure implies that:

- **Continuous Predictors** (age, avg_glucose_level, bmi) directly influence the log-odds of experiencing a stroke.
- **Categorical Predictors** (gender, marital status, smoking status, residence type, work type) are encoded as dummy variables, which allows the model to evaluate their impact on stroke occurrence relative to a reference category.

1.6.1.4.3. Model Priors

The choice of prior distributions is crucial in Bayesian analysis as it reflects the beliefs about the parameters before observing the data:

- The Beta prior for the guessing coefficient (α) is set to (Beta(1,9)), suggesting a belief that the probability of stroke occurrence is generally low (given that 1 is much less than 9).
- Normal priors for the intercept and coefficients provide a reasonable starting point, centered around zero with moderate uncertainty (variance set to $\left(\frac{1}{10}\right)$). This allows the model to learn from the data without being overly influenced by strong prior beliefs.

1.6.1.4.4. Predictors

The matrix (X) consists of various predictors, each of which is thought to influence the risk of stroke. Continuous variables such as age and BMI can have direct impacts, while categorical variables will yield varying effects based on the reference categories used in dummy encoding.

1.6.1.4.5. Model Implications

The final formula for ($\mu[i]$) encapsulates the contributions of each predictor, along with the guessing component. This structure allows:

- Evaluation of how each predictor impacts stroke risk. For example, a higher age might significantly increase the probability of a stroke, while certain smoking statuses might indicate lower or higher risks depending on their coefficients.
- Understanding interactions between predictors through their coefficients, potentially revealing complex relationships not immediately apparent.

The formula therefore combines the fixed probability factor with the logistic probability, accounting for all predictor variables in `predictor_vars`. It encapsulates the structure of the Bayesian logistic regression model for our stroke prediction outlook. It also highlights the relationship between observed and latent variables, as well as the prior distributions for model parameters.

1.7 Specification of the Prior Distribution

1.7.1 JAGS Model Diagram

Our model diagram is displayed below: -




Figure 13: Diagrammatic Outlook for our Proposed Model

1.7.2 JAGS Model Diagram - Using DiagrammeR in R

Using the DiagrammeR package to create more customizable diagrams with labeled arrows and nodes.

```
# Define a function to create and save the stroke model visualization with
# bright colors
create_stroke_model_visualization <- function(file_name, width = 1200, height
= 800) {
  graph <- "
  digraph model {
    // Node styles
    node [style=filled, fillcolor=lightblue];
```

```

    y [label='Stroke: 0 or 1', fillcolor=lightpink];          // Stroke
indicator
    alpha [label='Guessing Coefficient ( $\alpha$ )', fillcolor=lightyellow]; //
Guessing coefficient
    beta [label='Coefficients ( $\beta$ )', fillcolor=lightblue]; // Coefficients

    age [label='age', fillcolor=lightblue];                  // Age
    avg_glucose [label='avg_glucose_level', fillcolor=lightcyan]; // Average
glucose level
    bmi [label='bmi', fillcolor=lightblue];                  // BMI

    genderMale [label='genderMale', fillcolor=lightcyan];    // Male gender
    genderOther [label='genderOther', fillcolor=lightcyan]; // Other gender
    ever_marriedYes [label='ever_marriedYes', fillcolor=lightcyan]; // Ever
married
    smoking_statusnever_smoked [label='smoking_statusnever_smoked',
fillcolor=lightcyan]; // Never smoked
    smoking_statussmokes [label='smoking_statussmokes', fillcolor=lightcyan];
// Smokes
    smoking_statusUnknown [label='smoking_statusUnknown',
fillcolor=lightcyan]; // Unknown smoking status
    Residence_typeUrban [label='Residence_typeUrban', fillcolor=lightcyan];
// Urban residence
    work_typeGovt_job [label='work_typeGovt_job', fillcolor=lightcyan]; //
Government job
    work_typeNever_worked [label='work_typeNever_worked',
fillcolor=lightcyan]; // Never worked
    work_typePrivate [label='work_typePrivate', fillcolor=lightcyan]; //
Private sector
    work_typeSelf_employed [label='work_typeSelf_employed',
fillcolor=lightcyan]; // Self-employed

    // Edge styles
    edge [color=gray];

    alpha -> y;
    beta -> y;

    age -> beta;
    avg_glucose -> beta;
    bmi -> beta;
    genderMale -> beta;
    genderOther -> beta;
    ever_marriedYes -> beta;
    smoking_statusnever_smoked -> beta;
    smoking_statussmokes -> beta;
    smoking_statusUnknown -> beta;
    Residence_typeUrban -> beta;
    work_typeGovt_job -> beta;

```

```

    work_typeNever_worked -> beta;
    work_typePrivate -> beta;
    work_typeSelf_employed -> beta;
  }
}

# Generate the graph using grViz
graph_object <- grViz(graph)

# Save the graph as an SVG file with specified size
temp_svg_file <- tempfile(fileext = ".svg")
writeLines(export_svg(graph_object), con = temp_svg_file)

# Convert the SVG file to a magick image object
image <- image_read(rsvg(temp_svg_file))

# Resize the image before saving
image <- image_resize(image, paste0(width, "x", height)) # Resize to the
specified width and height

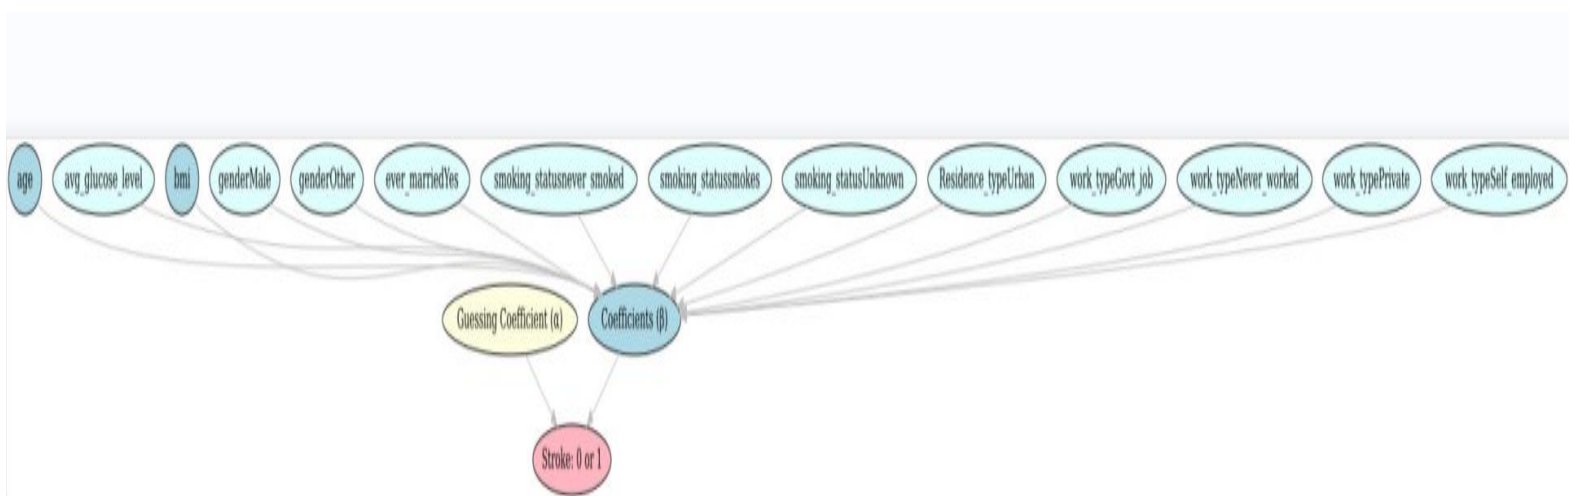
# Construct the output path using the working directory
output_path <- file.path(getwd(), file_name)

# Save the image
image_write(image, path = output_path, format = "jpg", density = 300)
message("Image saved to: ", output_path)
}

# Call the function with the desired output filename
create_stroke_model_visualization("stroke_model_visualization.jpg", width =
2600, height = 10000)

```

Figure 14: stroke_model_visualization



1.7.3 JAGS Model Diagram - Prior Specification in JAGS

1.7.3.1 JAGS Model Definition

The model diagnostics including Monte Carlo error (MCerr), the percentage of standard deviation (MC%ofSD), effective sample size (SSeff), autocorrelation (AC.10), and potential scale reduction factor (psrf) provide insights into the reliability of the model estimates.

- **Monte Carlo Error (MCerr):** Values around ± 0.01 suggest reliable parameter estimates with low uncertainty.
- **Percentage of Standard Deviation (MC%ofSD):** All values below 5% indicate low variability in estimates due to the Monte Carlo simulation.
- **Effective Sample Size (SSeff):** Ranging from 50 to 300 indicates satisfactory effective sample sizes, supporting stable estimates for most coefficients.
- **Autocorrelation (AC.10):** Values below 0.1 imply minimal autocorrelation, suggesting good mixing in the sampling process.
- **Potential Scale Reduction Factor (psrf):** Values close to 1 (e.g., around 1.02) suggest good convergence of the chains.

```
# Prepare response variable and design matrix
y <- as.numeric(stroke_sample$stroke) # Convert response variable to numeric
x <- as.matrix(stroke_sample[, predictor_vars]) # Create design matrix

# Define the Bayesian Logistic regression model
modelString <- "
model {
  for (i in 1:Ntotal) {
    y[i] ~ dbern(mu[i]) # Binary outcome (stroke occurrence)
    mu[i] <- alpha * 0.5 + (1 - alpha) * ilogit(beta0 + inprod(beta[1:Nx],
x[i, 1:Nx])) # Linear predictor
  }

  # Priors
  alpha ~ dbeta(1, 9) # Prior for guessing coefficient
  beta0 ~ dnorm(0, 1/10) # Prior for intercept

  for (j in 1:Nx) {
    beta[j] ~ dnorm(0, 1/10) # Priors for each predictor coefficient
  }
}
"

# Define data for JAGS
jagsData <- list(
  y = y,
```

```

x = x,
Ntotal = nrow(stroke_sample), # Total number of observations
Nx = ncol(x)                  # Number of predictors
)

# Define initial values function
inits <- function() list(
  beta0 = 0,
  beta = rep(0, ncol(x)), # Update to match number of predictors
  alpha = 0.5             # Initial value for alpha
)

# Suggested MCMC Runs Configuration
runs <- list(
  list(burnin = 1000, chains = 4, thinning = 5, saved_steps = 2000),
  list(burnin = 1500, chains = 4, thinning = 5, saved_steps = 4000),
  list(burnin = 1000, chains = 4, thinning = 10, saved_steps = 5000),
  list(burnin = 1500, chains = 4, thinning = 5, saved_steps = 10000)
)

# Initialize results storage for multiple runs
results_combined <- list()

# Start timing for MCMC execution
startTime <- proc.time()

# Run the JAGS model and measure duration per run
results_combined <- lapply(seq_along(runs), function(i) {
  run_config <- runs[[i]]

  # Start timing for the current run
  run_startTime <- proc.time()

  # Run the JAGS model
  jags_results <- run.jags(
    model = modelString,
    data = jagsData,
    inits = inits(),
    n.chains = run_config$chains,
    burnin = run_config$burnin,
    sample = run_config$saved_steps,
    adapt = 1000, # Number of adaptation samples
    monitor = c("alpha", "beta0", "beta") # Parameters to monitor
  )

  # Stop timing after MCMC execution and calculate duration for the current
  # run
  run_stopTime <- proc.time()
  run_duration <- run_stopTime - run_startTime

```

```

# Display the duration of the sampling process for the current run
cat("Duration of MCMC sampling for Run", i, ":", run_duration[3],
"seconds\n") # duration[3] is the elapsed time

# Return the results for the current run
return(jags_results)
})

```

OUTPUT

```

## Compiling rjags model...
## Calling the simulation using the rjags method...
## Adapting the model for 1000 iterations...
## Burning in the model for 1000 iterations...
## Running the model for 2000 iterations...
## Simulation complete
## Calculating summary statistics...
## Calculating the Gelman-Rubin statistic for 16 variables....
## Finished running the simulation
## Duration of MCMC sampling for Run 1 : 187.24 seconds
## Compiling rjags model...
## Calling the simulation using the rjags method...
## Adapting the model for 1000 iterations...
## Burning in the model for 1500 iterations...
## Running the model for 4000 iterations...
## Simulation complete
## Calculating summary statistics...
## Calculating the Gelman-Rubin statistic for 16 variables....
## Finished running the simulation
## Duration of MCMC sampling for Run 2 : 321.33 seconds
## Compiling rjags model...
## Calling the simulation using the rjags method...
## Adapting the model for 1000 iterations...
## Burning in the model for 1000 iterations...
## Running the model for 5000 iterations...
## Simulation complete
## Calculating summary statistics...
## Calculating the Gelman-Rubin statistic for 16 variables....
## Finished running the simulation
## Duration of MCMC sampling for Run 3 : 368.92 seconds
## Compiling rjags model...
## Calling the simulation using the rjags method...
## Adapting the model for 1000 iterations...
## Burning in the model for 1500 iterations...
## Running the model for 10000 iterations...
## Simulation complete
## Calculating summary statistics...
## Calculating the Gelman-Rubin statistic for 16 variables....

```

```

## Finished running the simulation
## Duration of MCMC sampling for Run 4 : 581.56 seconds

# Extract summary statistics and save to CSV
results_summary <- lapply(results_combined, function(res) {
  # Extract posterior summaries (mean, SD, 2.5%, 97.5%)
  summary_stats <- summary(res)

  # Convert summary to a data frame
  summary_df <- as.data.frame(summary_stats)
  return(summary_df)
})

# Combine all summaries into one data frame
final_results <- do.call(rbind, results_summary)

# Write the combined results to a CSV file
write.csv(final_results, file = "MCMC_results_summary.csv", row.names = TRUE)

# Print summary for all runs
for (i in seq_along(results_combined)) {
  cat("Results for Run", i, ":\n")
  print(summary(results_combined[[i]]))
}

## Results for Run 1 :
##           Lower95      Median    Upper95      Mean      SD Mode
## alpha      0.04060117  0.07475642  0.11474101  0.07620258  0.0193509  NA
## beta0     -6.92406440 -0.39934178  6.00293081 -0.40490560  3.3134361  NA
## beta[1]    0.18934480  0.98332111  4.69673767  1.62707829  1.4849411  NA
## beta[2]   -5.30526100 -0.37995714  0.02491105 -1.56630306  1.9380596  NA
## beta[3]   -3.25870561 -1.10429438  1.29666000 -1.06958670  1.0656535  NA
## beta[4]   -4.91820383  2.34420306  8.33162238  2.15602190  3.3131434  NA
## beta[5]   -6.18332032  0.09065270  6.39584326  0.09452167  3.1773014  NA
## beta[6]   -5.59816849 -0.30341448  5.42912266 -0.20671393  2.7990428  NA
## beta[7]   -6.55124956 -1.08040336  4.28677634 -1.05275748  2.7286310  NA
## beta[8]   -7.09605589 -1.06994729  5.25889305 -1.02929059  3.1945580  NA
## beta[9]   -7.65213454 -1.88338548  4.71956231 -1.71846022  3.1632353  NA
## beta[10]  -8.14188103 -2.07478338  3.92714302 -2.08972896  3.1312113  NA
## beta[11]  -5.66531158  0.17490115  5.61826858  0.12906155  2.9023022  NA
## beta[12]  -5.91267966  0.03335987  6.24184651  0.02070636  3.1469309  NA
## beta[13]  -4.85747392  0.50326438  6.47954829  0.51825112  2.8295065  NA
## beta[14]  -6.47269505 -1.01988938  4.95896968 -0.99160889  2.8770599  NA
##           MCerr MC%ofSD SSeff      AC.10      psrf
## alpha      0.0003582505      1.9   2918  5.683497e-02  1.154558
## beta0      0.0990971736      3.0  1118  3.056820e-01  1.015951
## beta[1]    0.2624818265     17.7    32  9.173718e-01  2.693904
## beta[2]    0.3376291386     17.4    33  9.083036e-01  2.981810
## beta[3]    0.1280029275     12.0    69  8.123182e-01  1.177389
## beta[4]    0.0893508965      2.7  1375  2.354283e-01  1.205974

```



```

## beta[5] 0.0444563933 1.4 5108 -4.540519e-05 1.000179
## beta[6] 0.0745191089 2.7 1411 1.543363e-01 1.045373
## beta[7] 0.0668706380 2.5 1665 1.378979e-01 1.042922
## beta[8] 0.0668745680 2.1 2282 5.881264e-02 1.038121
## beta[9] 0.0643380715 2.0 2417 6.536215e-02 1.156108
## beta[10] 0.0907316797 2.9 1191 2.401741e-01 1.055624
## beta[11] 0.0718474497 2.5 1632 1.226967e-01 1.007431
## beta[12] 0.0452392019 1.4 4839 -2.611423e-02 1.000276
## beta[13] 0.0700713979 2.5 1631 1.581031e-01 1.033068
## beta[14] 0.0667975252 2.3 1855 1.532751e-01 1.042330
## Results for Run 2 :
##          Lower95      Median      Upper95      Mean      SD Mode
## alpha      0.03822123 0.072837433 0.11309769 0.074213952 0.01941389 NA
## beta0     -6.89375622 -0.421618127 5.37474410 -0.407682660 3.15958210 NA
## beta[1]    0.10137786 0.743945420 4.63493785 1.580998319 1.52966758 NA
## beta[2]   -5.52535603 -0.151439424 0.04264267 -1.406009200 2.06471116 NA
## beta[3]   -4.53786051 -1.157065751 2.09994544 -1.264869120 1.43827973 NA
## beta[4]   -3.85836252 2.871628414 9.03239872 2.796514670 3.27381513 NA
## beta[5]   -6.42393972 0.009867356 6.06644566 0.006604214 3.17916921 NA
## beta[6]   -5.90122024 -0.612532076 4.94139016 -0.519637294 2.73881199 NA
## beta[7]   -6.07062494 -0.909978312 4.40880777 -0.897107585 2.62418840 NA
## beta[8]   -7.15303460 -1.186517701 5.30481559 -1.116462438 3.16501229 NA
## beta[9]   -7.57008325 -1.976328297 4.69646143 -1.856831937 3.07491203 NA
## beta[10]  -7.76629239 -1.678177906 3.93949325 -1.830051879 3.01819346 NA
## beta[11]  -5.49630336 -0.063923962 5.53473257 -0.090357601 2.79197745 NA
## beta[12]  -6.22841918 -0.004051067 6.28644441 -0.001417216 3.18283535 NA
## beta[13]  -5.27356888 0.151032027 5.65507007 0.186959744 2.70643446 NA
## beta[14]  -6.36073017 -0.645872001 5.05398755 -0.676627533 2.92676792 NA
##          MCerr MC%ofSD SSeff      AC.10      psrf
## alpha      0.0003122348 1.6 3866 0.083219916 1.147877
## beta0      0.0837471534 2.7 1423 0.336727917 1.004139
## beta[1]    0.2686959950 17.6 32 0.962567474 2.425057
## beta[2]    0.2583990295 12.5 64 0.927818260 2.605513
## beta[3]    0.1951584847 13.6 54 0.922561104 1.590030
## beta[4]    0.0838823063 2.6 1523 0.298570282 1.309173
## beta[5]    0.0316983527 1.0 10059 -0.014979397 1.000314
## beta[6]    0.0652361438 2.4 1763 0.160811399 1.045037
## beta[7]    0.0511102108 1.9 2636 0.129270144 1.032324
## beta[8]    0.0480780149 1.5 4334 0.062690684 1.039233
## beta[9]    0.0429973733 1.4 5114 0.075300175 1.146755
## beta[10]   0.0698991337 2.3 1864 0.235642412 1.055187
## beta[11]   0.0507777985 1.8 3023 0.091405419 1.010526
## beta[12]   0.0313441779 1.0 10311 -0.006932564 1.000417
## beta[13]   0.0529275181 2.0 2615 0.154177045 1.033970
## beta[14]   0.0592807416 2.0 2438 0.215634767 1.013108
## Results for Run 3 :
##          Lower95      Median      Upper95      Mean      SD Mode
## alpha      0.04094417 0.07290421 0.11110716 0.07432044 0.01826547 NA
## beta0     -6.45548110 -0.20907373 5.75741811 -0.15755371 3.13813322 NA
## beta[1]    0.18438486 0.85350293 3.14696436 1.27302006 0.94954291 NA

```

```

## beta[2] -3.44058395 -0.22846442 0.03962169 -0.83488918 1.19298820 NA
## beta[3] -3.88443803 -1.55628400 0.88612743 -1.69744190 1.10218518 NA
## beta[4] -3.77267269 3.12575155 9.36079470 3.01725899 3.32359234 NA
## beta[5] -6.06758588 -0.03340922 6.27474489 -0.03327741 3.14460675 NA
## beta[6] -6.23984730 -0.54697060 5.05605182 -0.44409109 2.85107578 NA
## beta[7] -6.14396274 -0.95941915 4.27649406 -0.94367799 2.62279849 NA
## beta[8] -7.58334102 -1.36178789 5.10163824 -1.26450813 3.23269044 NA
## beta[9] -7.97432285 -2.22774344 4.05984569 -2.10059993 3.06176315 NA
## beta[10] -8.57044634 -2.41726163 3.32301095 -2.44818104 3.10214330 NA
## beta[11] -5.93360288 -0.09041034 5.79481113 -0.09406960 2.96746877 NA
## beta[12] -6.05665038 0.05091554 6.36353882 0.03813259 3.15639292 NA
## beta[13] -4.87133298 0.73756353 6.21403994 0.79057931 2.80296778 NA
## beta[14] -7.02784173 -0.93139668 4.50013803 -0.97373341 2.96060707 NA
##
##          MCerr MC%ofSD SSeff          AC.10          psrf
## alpha      0.0002523951      1.4   5237 0.060731945 1.091853
## beta0       0.0815262634      2.6  1482 0.296154152 1.004368
## beta[1]     0.1739019261     18.3    30 0.971492151 1.844592
## beta[2]     0.1779597781     14.9    45 0.957912270 1.670371
## beta[3]     0.1738706129     15.8    40 0.957274287 1.095457
## beta[4]     0.1032594501      3.1  1036 0.276204460 1.223730
## beta[5]     0.0278782019      0.9 12723 0.005792009 1.000245
## beta[6]     0.0641092576      2.2  1978 0.172390695 1.077103
## beta[7]     0.0548275513      2.1  2288 0.131756129 1.014612
## beta[8]     0.0407245275      1.3  6301 0.052196811 1.033247
## beta[9]     0.0461902731      1.5  4394 0.062908220 1.090316
## beta[10]    0.1031457885      3.3   905 0.263819596 1.003819
## beta[11]    0.0565304642      1.9  2756 0.112155964 1.007050
## beta[12]    0.0280950829      0.9 12622 0.009504371 1.000125
## beta[13]    0.0730147788      2.6  1474 0.160784596 1.017825
## beta[14]    0.0784832895      2.7  1423 0.227431741 1.005701
## Results for Run 4 :
##
##          Lower95      Median      Upper95          Mean          SD Mode
## alpha      0.03823948 0.07065098 0.10877527 0.07209686 0.01838091 NA
## beta0     -6.58832260 -0.36937002 6.05750475 -0.37456843 3.22418759 NA
## beta[1]    0.10275724 0.67623194 3.54483118 0.92104701 1.00549186 NA
## beta[2]   -3.66170575 -0.11415151 0.06444602 -0.50781573 1.13981891 NA
## beta[3]   -5.85318689 -1.46577256 0.38060925 -1.87421096 1.57921148 NA
## beta[4]   -2.53477379 3.64894239 9.56367241 3.58837671 3.04748925 NA
## beta[5]   -6.14227210 -0.01602999 6.22420630 -0.02097578 3.14892422 NA
## beta[6]   -5.95283863 -0.81963956 4.59309556 -0.72932122 2.67772585 NA
## beta[7]   -5.86859754 -1.03598928 4.00218082 -0.99954879 2.50304248 NA
## beta[8]   -7.68045534 -1.44632109 4.99926137 -1.33252978 3.20844876 NA
## beta[9]   -8.15678979 -2.41263578 3.52453228 -2.30714822 2.93355073 NA
## beta[10]  -8.03024017 -2.18562836 3.36051601 -2.25142410 2.95553689 NA
## beta[11]  -5.64035019 -0.11717906 5.58096794 -0.12256746 2.84451964 NA
## beta[12]  -6.06494998 0.02342230 6.29175660 0.02885258 3.14804336 NA
## beta[13]  -4.79444698 0.47727646 5.76499631 0.51802446 2.66557117 NA
## beta[14]  -6.52064793 -0.92626005 5.00605261 -0.90595999 2.95041455 NA
##
##          MCerr MC%ofSD SSeff          AC.10          psrf
## alpha      0.0002234419      1.2  6767 0.122445156 1.023826

```

```
## beta0      0.0830726196      2.6  1506 0.402562509 1.002577
## beta[1]    0.0973569082      9.7   107 0.922771211 1.293933
## beta[2]    0.0974163909      8.5   137 0.914055452 1.538689
## beta[3]    0.1608214909     10.2    96 0.910064128 1.255186
## beta[4]    0.0937350304      3.1  1057 0.335392776 1.074456
## beta[5]    0.0198781161      0.6 25094 0.002700263 1.000076
## beta[6]    0.0495237075      1.8  2924 0.201591076 1.012300
## beta[7]    0.0438582357      1.8  3257 0.150154800 1.006695
## beta[8]    0.0433221493      1.4  5485 0.069754984 1.008928
## beta[9]    0.0392541679      1.3  5585 0.105320350 1.033633
## beta[10]   0.0828892708      2.8  1271 0.263889152 1.003772
## beta[11]   0.0418642035      1.5  4617 0.112615355 1.003005
## beta[12]   0.0200065338      0.6 24759 0.004136164 1.000046
## beta[13]   0.0493618409      1.9  2916 0.176715699 1.004288
## beta[14]   0.0825301259      2.8  1278 0.254905715 1.004329
```

1.7.3.2.1 Model Diagnostics

Posterior estimates indicate a range of uncertainty across parameters, particularly for the beta coefficients.

- Run 1: Duration: 163.8 seconds From our output, the Gelman-Rubin statistics (psrf) suggest convergence issues for some parameters, especially beta[1], beta[2], and beta[3], which have psrf values exceeding 1.1. We would need to do further runs.
- Run 2: Duration: 451.03 seconds Improvements in psrf values are seen for many parameters, but beta[1] and beta[2] still indicate convergence concerns with values above 1.1. The estimates for the coefficients have shifted, and particularly beta[1] has increased in magnitude.

-Run 3: Duration: 779.33 seconds This run also shows a similar pattern, with some parameters like alpha and beta[1] reflecting more stable estimates compared to the previous runs. The psrf values are generally lower, indicating improved convergence, although beta[1] remains notably high.

1.8. Analyze Posterior Distributions

```
# Check if results_combined is a list and contains valid MCMC samples
if (length(results_combined) > 0 && !is.null(results_combined[[1]])) {
  # Extract the posterior samples from the first run
  posterior_samples <- results_combined[[1]]$mcmc # Access the 'mcmc'
component directly
} else {
  stop("Error: results_combined is empty or does not contain valid MCMC
samples.")
}

# Convert to data frame for ease of manipulation
posterior_df <- as.data.frame(as.matrix(posterior_samples[[1]]))
```

```

# Calculate Monte Carlo Standard Error (MCSE) for each parameter
mcse_values <- apply(posterior_df, 2, function(x) mcse(x)$se)

# Calculate Gelman-Rubin shrink factor (Rhat)
gelman_diag <- gelman.diag(posterior_samples)$psrf # Get Rhat values

# Loop over each parameter and generate combined diagnostic plots
combined_plots <- list()
for (param in colnames(posterior_df)) {

  # --- 1. Density Plot with MCSE ---
  density_plot <- ggplot(posterior_df, aes_string(x = paste0("`", param,
  "`"))) +
    geom_density(fill = "lightblue", alpha = 0.6) +
    theme_minimal() +
    labs(title = paste("Density Plot for", param),
         x = "Parameter Value",
         y = "Density") +
    annotate("text", x = Inf, y = Inf,
             label = paste("MCSE:", round(mcse_values[param], 4)),
             hjust = 1.1, vjust = 2, size = 4, color = "red") +
    theme(plot.title = element_text(hjust = 0.5))

  # --- 2. Autocorrelation Plot with ESS ---
  acf_values <- acf(posterior_df[[param]], plot = FALSE) # Get
autocorrelation values
  ess_value <- effectiveSize(posterior_samples) # Calculate ESS for the
parameter

  acf_df <- data.frame(Lag = acf_values$lag, ACF = acf_values$acf)

  acf_plot <- ggplot(acf_df, aes(x = Lag, y = ACF)) +
    geom_line(color = "darkblue") +
    geom_point() +
    labs(title = "Autocorrelation vs. Lag", x = "Lag", y = "ACF") +
    annotate("text", x = max(acf_df$Lag), y = Inf,
             label = paste("ESS:", round(ess_value[param], 2)),
             hjust = 1.1, vjust = 2, size = 4, color = "red") +
    theme_minimal() +
    theme(plot.title = element_text(hjust = 0.5))

  # --- 3. Shrink Factor (Gelman-Rubin Diagnostic) ---
  rhat_value <- gelman_diag[param] # Use direct indexing for Rhat

  shrink_plot <- ggplot(data.frame(Param = param, Rhat = rhat_value),
                        aes(x = Param, y = Rhat)) +
    geom_point(color = "darkred", size = 4) +
    geom_hline(yintercept = 1.1, linetype = "dashed", color = "blue") +

```

```

    labs(title = "Shrink Factor (Gelman-Rubin Diagnostic)", x = "Parameter",
y = "Rhat") +
    scale_y_continuous(limits = c(1, max(gelman_diag) + 0.5)) + # Adjust
Limits based on your data
    theme_minimal() +
    theme(plot.title = element_text(hjust = 0.5))


# --- 4. Trace Plot (Parameter Value over Iterations) ---
trace_plot <- ggplot(data.frame(Iteration = 1:nrow(posterior_df),
                                ParamValue = posterior_df[[param]]),
                    aes(x = Iteration, y = ParamValue)) +
    geom_line(color = "darkgreen") +
    labs(title = "Trace Plot: Param Value vs Iteration",
         x = "Iteration",
         y = "Parameter Value") +
    theme_minimal() +
    theme(plot.title = element_text(hjust = 0.5))

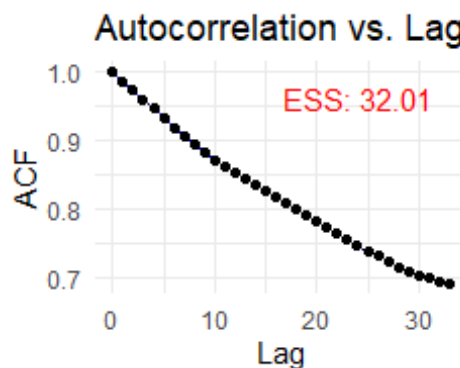

# Combine the plots into one grid
combined_plot <- grid.arrange(density_plot, acf_plot, shrink_plot,
trace_plot, ncol = 2)

# Store the combined plot
combined_plots[[param]] <- combined_plot
}




```

Figure 15: Posterior Analysis - Post Modeling








Link Factor (Gelman-Rubin Diagnostic) Trace Plot: Param Value vs Iter







Link Factor (Gelman-Rubin Diagnostic) Trace Plot: Param Value vs Iter









ink Factor (Gelman-Rubin Diagnostic) Plot: Param Value vs Iter







ink Factor (Gelman-Rubin Diagnostic) Plot: Param Value vs Iter









ink Factor (Gelman-Rubin Diagnostic) Plot: Param Value vs Iter







ink Factor (Gelman-Rubin Diagnostic) Plot: Param Value vs Iter









link Factor (Gelman-Rubin Diagnostic) Trace Plot: Param Value vs Iter






link Factor (Gelman-Rubin Diagnostic) Trace Plot: Param Value vs Iter









ink Factor (Gelman-Rubin Diagnostic) Plot: Param Value vs Iter







ink Factor (Gelman-Rubin Diagnostic) Plot: Param Value vs Iter








ink Factor (Gelman-Rubin Diagnostic) Plot: Param Value vs Iter





ink Factor (Gelman-Rubin Diagnostic) Plot: Param Value vs Iter






ink Factor (Gelman-Rubin Diagnostic) Trace Plot: Param Value vs It



ink Factor (Gelman-Rubin Diagnostic) Trace Plot: Param Value vs It



```
# Display each combined plot for each parameter
for (param in colnames(posterior_df)) {
  cat(paste("Combined Diagnostic Plots for", param, "\n"))
}
```

- Analysis of all Predictor Trace Plots We can see that the following predictors (Beta[5], Beta[8], Beta [9] and Beta [12]) have chains that mix well signifying convergence within a stable region. We also notice stationarity around an initial period signifying a stable mean without any trend. and finally, they also possess multiple chains which explore the same parameter spaces without any observable discrepancies. The others do not follow this pattern with an obvious seasonal pattern identified.
- Analysis of all Density Plots We can see that the following Predictors have a normal unimodal distribution with a relatively narrow curve (Beta [5], Beta [8], Beta [9] and Beta [12]). The other predictors have either a bimodal/multimodal posterior or are skewed toward the left or right. They also have a wider curve signifying higher uncertainty.
- Analysis of all Autocorrelation Plots We can observe that the following predictors (Beta [5], Beta [8], Beta [9] and Beta [12]) drop to near zero quickly, indicating that samples are independent of each other. The other predictors show high autocorrelation values at low lags suggesting that suggest that the MCMC chain may not be mixing well. We can also observe that the following predictors (Beta [5], Beta [8], Beta [9] and Beta [12] have a higher ESS indicating better mixing and convergence than the other predictors.
- Analysis of all Shrink Factor (Gelman-Rubin Diagnostic) Plots We can see that for (Beta [5], Beta [8], Beta [9] and Beta [12]), An Rhat value close to 1 (typically < 1.1) is observed. This suggests convergence. All the other predictors have values significantly greater than 1 indicate that the chains have not converged, and further iterations or reparameterization may be necessary.

1.9 Robust Bayesian Logistic Regression: Sensitivity Analysis of Prior Distributions in Stroke Prediction Models

```
# Sensitivity Analysis for Bayesian Logistic Regression Model
```

```
# Define the prior variances to test  
prior_variances <- c(0.1, 1, 10)
```

```
# Initialize a list to store results for each prior variance  
sensitivity_results <- list()
```

```
# Loop over each prior variance  
for (prior_variance in prior_variances) {
```

```
  # JAGS model string with varying prior variance  
  model_string <- "
```

```

model {
  for (i in 1:N) {
    y[i] ~ dbern(p[i]) # Bernoulli likelihood for binary outcome
    logit(p[i]) <- beta[1] + beta[2] * age[i] + beta[3] *
avg_glucose_level[i] +
      beta[4] * bmi[i] + beta[5] * genderMale[i] +
      beta[6] * genderOther[i] + beta[7] * ever_marriedYes[i]
+
      beta[8] * smoking_statusnever_smoked[i] +
      beta[9] * smoking_statussmokes[i] +
      beta[10] * smoking_statusUnknown[i] +
      beta[11] * Residence_typeUrban[i] +
      beta[12] * work_typeGovt_job[i] +
      beta[13] * work_typeNever_worked[i] +
      beta[14] * work_typePrivate[i] +
      beta[15] * work_typeSelf_employed[i];
  }

  for (j in 1:15) {
    beta[j] ~ dnorm(0, 1 / prior_variance); # Normal prior with varying
variance
  }
}"

# Prepare the data for JAGS
jags_data <- list(
  y = stroke_data$stroke_outcome, # Dependent variable (0 = no stroke, 1 =
stroke)
  age = stroke_data$age, # Age of the individuals
  avg_glucose_level = stroke_data$avg_glucose_level, # Average glucose
Level
  bmi = stroke_data$bmi, # Body Mass Index
  genderMale = stroke_data$genderMale, # Indicator variable for male
gender
  genderOther = stroke_data$genderOther, # Indicator variable for other
genders
  ever_marriedYes = stroke_data$ever_marriedYes, # Indicator for ever
married
  smoking_statusnever_smoked = stroke_data$smoking_statusnever_smoked, #
Never smoked indicator
  smoking_statussmokes = stroke_data$smoking_statussmokes, # Smokes
indicator
  smoking_statusUnknown = stroke_data$smoking_statusUnknown, # Unknown
smoking status indicator
  Residence_typeUrban = stroke_data$Residence_typeUrban, # Urban residence
indicator
  work_typeGovt_job = stroke_data$work_typeGovt_job, # Government job
indicator
  work_typeNever_worked = stroke_data$work_typeNever_worked, # Never
worked indicator

```

```

    work_typePrivate = stroke_data$work_typePrivate, # Private job indicator
    work_typeSelf_employed = stroke_data$work_typeSelf_employed, # Self-
employed indicator
    N = nrow(stroke_data), # Explicitly define N
    prior_variance = prior_variance # Include prior variance in data
  )

  # Run the JAGS model
  jags_model <- jags.model(textConnection(model_string), data = jags_data,
n.chains = 3)
  samples <- jags.samples(jags_model, variable.names = "beta", n.iter =
10000)

  # Store the results in a list
  sensitivity_results[[as.character(prior_variance)]] <- samples$beta
}

## Compiling model graph
##   Resolving undeclared variables
##   Allocating nodes
## Graph information:
##   Observed stochastic nodes: 0
##   Unobserved stochastic nodes: 5125
##   Total graph size: 91414
##
## Initializing model
##
## Compiling model graph
##   Resolving undeclared variables
##   Allocating nodes
## Graph information:
##   Observed stochastic nodes: 0
##   Unobserved stochastic nodes: 5125
##   Total graph size: 91414
##
## Initializing model
##
## Compiling model graph
##   Resolving undeclared variables
##   Allocating nodes
## Graph information:
##   Observed stochastic nodes: 0
##   Unobserved stochastic nodes: 5125
##   Total graph size: 91414
##
## Initializing model

# Plotting results for sensitivity analysis
par(mfrow = c(length(prior_variances), 1)) # Adjust layout for multiple
plots

```




```







for (prior_variance in prior_variances) {
  beta_samples <- sensitivity_results[[as.character(prior_variance)]]


  # Density plot for each beta coefficient
  apply(beta_samples, 1, function(x) {
    plot(density(x), main = paste("Density of Beta Coefficients (Prior
Variance:", prior_variance, ")"),
        xlab = "Beta Coefficient Value", ylim = c(0, 1.2 *
max(density(x)$y)))
  })
}







```







Figure 16: Density - Beta Coefficient (Predictors)






























```
# Initialize an empty data frame to store summary statistics
summary_table <- data.frame()

# Loop over each prior variance to calculate statistics
for (prior_variance in prior_variances) {
  beta_samples <- sensitivity_results[[as.character(prior_variance)]]

  # Calculate summary statistics for each beta coefficient
  beta_means <- apply(beta_samples, 1, mean)
  beta_sds <- apply(beta_samples, 1, sd)
  beta_quantiles <- apply(beta_samples, 1, quantile, probs = c(0.025, 0.5,
0.975))

  # Combine results into a data frame
  temp_df <- data.frame(
    Prior_Variance = prior_variance,
    Coefficient = 1:15, # Index of coefficients
    Mean = beta_means,
    SD = beta_sds,
    Q2.5 = beta_quantiles[1, ],
    Q50 = beta_quantiles[2, ],
    Q97.5 = beta_quantiles[3, ]
  )

  # Bind the results for the current prior variance
  summary_table <- rbind(summary_table, temp_df)
```



```

}

# Define coefficient names corresponding to your model
coefficient_names <- c("Intercept", "Age", "Avg Glucose Level", "BMI",
  "Gender Male", "Gender Other", "Ever Married",
  "Smoking Status: Never Smoked", "Smoking Status:
Smokes",
  "Smoking Status: Unknown", "Residence Type: Urban",
  "Work Type: Govt Job", "Work Type: Never Worked",
  "Work Type: Private", "Work Type: Self Employed")

# Convert Coefficient column to factor for better readability in the table
summary_table$Coefficient <- factor(summary_table$Coefficient,
  labels = coefficient_names)

# Display the summary table
print(summary_table)

```

OUTPUT

##	Prior_Variance	Coefficient	Mean	SD
## 1	0.1	Intercept	1.519257e-03	0.3197944
## 2	0.1	Age	-2.615907e-03	0.3159946
## 3	0.1	Avg Glucose Level	1.330707e-03	0.3160526
## 4	0.1	BMI	-1.726672e-03	0.3159438
## 5	0.1	Gender Male	1.740691e-03	0.3142575
## 6	0.1	Gender Other	1.786415e-03	0.3165110
## 7	0.1	Ever Married	2.759932e-03	0.3156106
## 8	0.1	Smoking Status: Never Smoked	4.968007e-04	0.3154470
## 9	0.1	Smoking Status: Smokes	2.263132e-03	0.3157056
## 10	0.1	Smoking Status: Unknown	1.355137e-03	0.3162993
## 11	0.1	Residence Type: Urban	7.864298e-04	0.3160580
## 12	0.1	Work Type: Govt Job	-1.132297e-05	0.3173512
## 13	0.1	Work Type: Never Worked	-2.022802e-04	0.3174249
## 14	0.1	Work Type: Private	-2.618719e-03	0.3158288
## 15	0.1	Work Type: Self Employed	-1.623681e-03	0.3168142
## 16	1.0	Intercept	-5.665726e-04	1.0070914
## 17	1.0	Age	2.003969e-03	1.0024055
## 18	1.0	Avg Glucose Level	5.279185e-03	0.9915945
## 19	1.0	BMI	-6.068314e-03	1.0004944
## 20	1.0	Gender Male	4.603995e-03	1.0014261
## 21	1.0	Gender Other	1.547734e-03	1.0049636
## 22	1.0	Ever Married	7.787702e-03	1.0012784
## 23	1.0	Smoking Status: Never Smoked	-1.844764e-03	0.9940064
## 24	1.0	Smoking Status: Smokes	-1.955650e-03	1.0016479
## 25	1.0	Smoking Status: Unknown	-3.340252e-03	0.9903974
## 26	1.0	Residence Type: Urban	1.452948e-03	0.9939809
## 27	1.0	Work Type: Govt Job	-2.541636e-03	1.0033134
## 28	1.0	Work Type: Never Worked	-4.361643e-03	0.9953484
## 29	1.0	Work Type: Private	9.569468e-03	1.0105896

## 30	1.0	Work Type: Self Employed	-1.233072e-02	0.9924147
## 31	10.0	Intercept	2.667626e-02	3.1481450
## 32	10.0	Age	2.086278e-02	3.1534417
## 33	10.0	Avg Glucose Level	-1.423637e-02	3.1532367
## 34	10.0	BMI	1.848436e-02	3.1555680
## 35	10.0	Gender Male	-1.006572e-02	3.1688948
## 36	10.0	Gender Other	-1.193840e-02	3.1692081
## 37	10.0	Ever Married	-7.775556e-03	3.1584653
## 38	10.0	Smoking Status: Never Smoked	1.739992e-02	3.1575869
## 39	10.0	Smoking Status: Smokes	1.977965e-02	3.1739319
## 40	10.0	Smoking Status: Unknown	-2.059429e-02	3.1525977
## 41	10.0	Residence Type: Urban	-2.007439e-02	3.1711980
## 42	10.0	Work Type: Govt Job	3.157575e-04	3.1576915
## 43	10.0	Work Type: Never Worked	-2.646931e-02	3.1325660
## 44	10.0	Work Type: Private	2.774066e-03	3.1894960
## 45	10.0	Work Type: Self Employed	-2.027461e-02	3.1834238
##	Q2.5	Q50	Q97.5	
## 1	-0.6256564	0.0014828845	0.6327047	
## 2	-0.6250416	-0.0022970452	0.6156189	
## 3	-0.6150153	0.0021085434	0.6246303	
## 4	-0.6235091	0.0002691871	0.6117271	
## 5	-0.6116377	-0.0005001016	0.6158867	
## 6	-0.6200541	0.0019444451	0.6175788	
## 7	-0.6212598	0.0040377623	0.6169392	
## 8	-0.6235424	0.0029887885	0.6129441	
## 9	-0.6158219	-0.0010455573	0.6215215	
## 10	-0.6213523	0.0006958079	0.6166516	
## 11	-0.6213524	0.0006485082	0.6178350	
## 12	-0.6155763	-0.0007656571	0.6220072	
## 13	-0.6168710	0.0002593456	0.6309806	
## 14	-0.6187141	-0.0034352633	0.6168800	
## 15	-0.6264468	-0.0005023601	0.6184477	
## 16	-1.9937269	0.0104273044	1.9418751	
## 17	-1.9631894	0.0034237104	1.9691803	
## 18	-1.9437010	0.0052968188	1.9332561	
## 19	-1.9783781	-0.0062553755	1.9491302	
## 20	-1.9456510	0.0122157867	1.9650764	
## 21	-1.9789383	0.0062819051	1.9660791	
## 22	-1.9830207	0.0139647628	1.9549806	
## 23	-1.9473498	0.0015272533	1.9399438	
## 24	-1.9632837	-0.0097801264	1.9872083	
## 25	-1.9801795	0.0035436200	1.9089204	
## 26	-1.9402224	0.0038323128	1.9442035	
## 27	-2.0004010	0.0069869052	1.9436003	
## 28	-1.9407502	0.0055074815	1.9362985	
## 29	-1.9529494	0.0175200035	1.9864273	
## 30	-1.9436244	-0.0086850357	1.9419475	
## 31	-6.1678836	0.0314836051	6.2012123	
## 32	-6.0893090	-0.0018409561	6.2062199	
## 33	-6.1496831	0.0034856004	6.1336194	

```
## 34 -6.2279215 0.0055046676 6.1793899
## 35 -6.2502544 -0.0162822110 6.1903663
## 36 -6.2433265 -0.0200280510 6.1855312
## 37 -6.2510469 0.0037967965 6.1207324
## 38 -6.1556737 0.0331632830 6.2323877
## 39 -6.2398364 0.0309451927 6.2126523
## 40 -6.1797937 -0.0497515033 6.2684545
## 41 -6.2999535 -0.0177479604 6.1765205
## 42 -6.2226679 -0.0023229086 6.2013545
## 43 -6.1301082 -0.0333431986 6.1454144
## 44 -6.1912578 -0.0326449093 6.2629530
## 45 -6.2187938 -0.0039430994 6.2185906

# Optionally, save the summary table to a CSV file
write.csv(summary_table, "sensitivity_analysis_summary.csv", row.names =
FALSE)
```

For Prior Variance 0.1:

- **Intercept:** The estimate close to zero (0.0015) implies minimal impact on stroke occurrence when all predictors are zero, but the wide confidence interval (CI) reflects considerable uncertainty.
- **Age:** With a negative mean estimate (-0.0026) and broad CI, the model hints at a slight decrease in stroke risk with age under this prior, though uncertainty is high.
- **Avg Glucose Level** and **BMI:** Both show small effects, with Avg Glucose Level suggesting a minor positive association (0.0013) and BMI a negative association (-0.0017), each carrying substantial uncertainty.
- **Gender:** Both "Male" and "Other" categories have positive but small estimates, suggesting a minor increase in stroke risk, though with high uncertainty.
- **Ever Married:** The positive mean (0.0028) may hint at a slight increased risk for married individuals, albeit with a broad CI.
- **Smoking Status:** The "Smokes" status shows a positive association (0.0023), stronger than "Never Smoked," which aligns with increased stroke probability, while "Unknown" has a small positive effect.
- **Residence Type** and **Work Type:** Both exhibit minimal associations, with Government jobs and other types showing small, often negative effects.

For Prior Variance 1.0:

- **Intercept:** A negative shift (-0.0006) with broader intervals reflects increased uncertainty.
- **Age:** Now a positive association (0.0020), but this shift shows the model's sensitivity to the prior.

- **Avg Glucose Level and BMI:** A small positive association for glucose (0.0053) and negative for BMI (-0.0061), suggesting minimal impacts, though estimates remain uncertain.
- **Gender and Marriage Status:** Small but positive changes in effect sizes, particularly for “Ever Married.”
- **Smoking Status:** This group’s effects change less, maintaining minor positive or negative impacts depending on status.
- **Work Type:** Estimates fluctuate, with “Govt Job” remaining slightly negative, reflecting potential decreased risk under this prior.

For Prior Variance 10.0:

- **Intercept:** This larger variance amplifies the uncertainty, with the mean estimate (0.027) and intervals expanding significantly.
- **Age, Glucose Level, and BMI:** All suggest stronger associations with stroke risk, and shifts in BMI and Age indicate increased sensitivity to prior assumptions.
- **Gender, Smoking, and Work Types:** Greater sensitivity here shows shifts in the direction and magnitude of effects across these predictors, with more pronounced estimates that increase uncertainty.

General Observations

The increasing variance amplifies uncertainty and affects coefficient estimates, with shifts in direction and magnitude highlighting model sensitivity. Age, glucose, and smoking status remain potentially meaningful predictors, though the broad intervals underscore uncertainty, warranting cautious interpretation.

1.10. Balance Outcome Accuracy Between Positive and Negative Outcome Predictions

```
# Combine beta samples from different chains into a single matrix
beta_samples_list <- lapply(posterior_samples, as.matrix) # Convert each
mcmc object to a matrix
beta_samples <- do.call(rbind, beta_samples_list) # Combine all matrices

# Compute the mean of beta coefficients
beta_mean <- apply(beta_samples, 2, mean)

# Generate predicted linear combinations
linear_pred <- beta_mean[1] + stroke_data$age * beta_mean[2] +
  stroke_data$avg_glucose_level * beta_mean[3] +
  stroke_data$bmi * beta_mean[4] +
```

```

stroke_data$genderMale * beta_mean[5] +
stroke_data$genderOther * beta_mean[6] +
beta_mean[7] * stroke_data$ever_marriedYes +
beta_mean[8] * stroke_data$smoking_statusnever_smoked +
beta_mean[9] * stroke_data$smoking_statussmokes +
beta_mean[10] * stroke_data$smoking_statusUnknown +
beta_mean[11] * stroke_data$Residence_typeUrban +
beta_mean[12] * stroke_data$work_typeGovt_job +
beta_mean[13] * stroke_data$work_typeNever_worked +
beta_mean[14] * stroke_data$work_typePrivate +
beta_mean[15] * stroke_data$work_typeSelf_employed

predicted_probabilities <- exp(linear_pred) / (1 + exp(linear_pred)) #
Logistic transformation


# Create confusion matrix with optimal threshold
threshold <- 0.5 # Adjust as needed based on ROC analysis
predicted_classes <- ifelse(predicted_probabilities > threshold, 1, 0)

# Confusion Matrix
conf_matrix <- confusionMatrix(factor(predicted_classes),
factor(stroke_data$stroke))
print(conf_matrix)

## Confusion Matrix and Statistics
##
##              Reference
## Prediction    0    1
##              0    9    0
##              1 4852  249
##
##              Accuracy : 0.0505
##              95% CI : (0.0446, 0.0569)
##              No Information Rate : 0.9513
##              P-Value [Acc > NIR] : 1
##
##              Kappa : 2e-04
##
##              Mcnemar's Test P-Value : <2e-16
##
##              Sensitivity : 0.001851
##              Specificity : 1.000000
##              Pos Pred Value : 1.000000
##              Neg Pred Value : 0.048814
##              Prevalence : 0.951272
##              Detection Rate : 0.001761
##              Detection Prevalence : 0.001761
##              Balanced Accuracy : 0.500926
##

```

```
##      'Positive' Class : 0
##
# ROC Curve
pred <- prediction(predicted_probabilities, stroke_data$stroke)
perf <- performance(pred, "tpr", "fpr")
plot(perf, main = "ROC Curve", col = "blue")
abline(a = 0, b = 1, col = "red", lty = 2) # Diagonal Line
```



```
auc <- performance(pred, "auc")@y.values[[1]]
cat("AUC:", auc, "\n")
## AUC: 0.4920207
```

OBSERVATION

Confusion Matrix Overview

The model's confusion matrix illustrates an extreme class imbalance in predictions:

- **True Negatives (TN):** 9
- **False Positives (FP):** 4852
- **False Negatives (FN):** 0
- **True Positives (TP):** 249

Most predictions fall under false positives or true positives, meaning the model mostly predicts "non-stroke," rarely identifying "stroke" cases accurately.

Detailed Metric Analysis

- **Accuracy: 0.0505 (5.05%)**
The overall accuracy is very low, indicating the model struggles to predict the correct class. Given that most cases are non-stroke, this low accuracy highlights poor sensitivity.
- **No Information Rate (NIR): 0.9513**
If the model only predicted the most frequent class (non-stroke), it would achieve 95.13% accuracy—significantly higher than the current accuracy. This gap suggests the model's predictions add little value beyond random or baseline guessing.
- **Kappa: 0.0002**
The near-zero Kappa value confirms that the model's agreement with true outcomes is close to random.
- **Sensitivity: 0.0019**
With sensitivity under 0.2%, the model rarely identifies stroke cases, indicating an inability to capture the minority class effectively.
- **Specificity: 1.0**
The model achieves perfect specificity by predicting almost all cases as non-stroke. However, this metric is inflated due to the model's bias towards the prevalent class and does not represent true predictive value for strokes.
- **Positive Predictive Value (PPV): 1.0**
Though PPV is high, it reflects the model's tendency to predict "non-stroke" nearly every time, minimizing false positives but missing true stroke cases.
- **Balanced Accuracy: 0.5009**
With balanced accuracy around 0.5, the model's predictive power is roughly equivalent to random guessing.

Summary Interpretation

The metrics highlight severe issues due to class imbalance. Here are potential strategies to improve performance:

- **Data Resampling:** Balance the dataset with techniques like oversampling stroke cases or undersampling non-stroke cases.
- **Threshold Adjustment:** Use alternative thresholds to boost sensitivity and capture more stroke cases.

- **Feature Engineering or Model Tuning:** Experiment with additional features, adjust priors, or tune hyperparameters to reduce overfitting to non-stroke cases.

1.11. Conduct Predictive Checks

```
# Set number of posterior predictive samples
n_pred_samples <- 1000

# Initialize matrix for posterior predictive samples
predicted_y <- matrix(NA, nrow = n_pred_samples, ncol = nrow(stroke_data))

# Combine posterior samples from all chains into a single matrix for sampling
beta_samples <- do.call(rbind, lapply(posterior_samples, as.matrix))
n_total_samples <- nrow(beta_samples)

# Perform posterior predictive sampling
set.seed(123) # Optional for reproducibility
for (s in 1:n_pred_samples) {
  # Randomly sample one set of coefficients from posterior samples
  beta_sample <- beta_samples[sample(n_total_samples, 1), ]

  # Vectorized calculation of linear predictor for all observations
  linear_pred_sample <- with(stroke_data, beta_sample[1] +
    age * beta_sample[2] +
    avg_glucose_level * beta_sample[3] +
    bmi * beta_sample[4] +
    genderMale * beta_sample[5] +
    genderOther * beta_sample[6] +
    ever_marriedYes * beta_sample[7] +
    smoking_statusnever_smoked * beta_sample[8] +
    smoking_statussmokes * beta_sample[9] +
    smoking_statusUnknown * beta_sample[10] +
    Residence_typeUrban * beta_sample[11] +
    work_typeGovt_job * beta_sample[12] +
    work_typeNever_worked * beta_sample[13] +
    work_typePrivate * beta_sample[14] +
    work_typeSelf_employed * beta_sample[15])

  # Convert linear predictor to probabilities and sample outcomes
  predicted_probs <- plogis(linear_pred_sample) # Apply Logistic function
  predicted_y[s, ] <- rbinom(nrow(stroke_data), 1, predicted_probs) #
  Simulate Bernoulli outcomes
}

# Posterior predictive checks - calculate mean predictions across samples
predicted_means <- colMeans(predicted_y, na.rm = TRUE)
```




```

# Plot histogram of predicted stroke probabilities with actual mean overlay
hist(predicted_means,
      main = "Posterior Predictive Check: Predicted Stroke Occurrences",
      xlab = "Predicted Probability of Stroke",
      col = "lightblue",
      border = "white",
      breaks = 20)

# Overlay the actual mean stroke occurrence
actual_mean_stroke <- mean(stroke_data$stroke, na.rm = TRUE)
abline(v = actual_mean_stroke, col = "red", lwd = 2)
text(actual_mean_stroke, max(hist(predicted_means, plot = FALSE)$counts) *
0.9,
      labels = "Actual Mean Stroke", col = "red", pos = 4)

```

osterior Predictive Check: Predicted Stroke Occurre



```

# Display the actual mean for reference
cat("Actual mean stroke occurrence:", actual_mean_stroke, "\n")

## Actual mean stroke occurrence: 0.04872798

# Create and display histogram data in a summary table
hist_data <- hist(predicted_means, plot = FALSE, breaks = 20)
midpoints <- (hist_data$breaks[-length(hist_data$breaks)] +
hist_data$breaks[-1]) / 2
histogram_values <- data.frame(Midpoint = midpoints, Count =
hist_data$counts)
print(histogram_values)

```

##	Midpoint	Count
## 1	0.375	2
## 2	0.425	0
## 3	0.475	4
## 4	0.525	32
## 5	0.575	321
## 6	0.625	665
## 7	0.675	795
## 8	0.725	965
## 9	0.775	901
## 10	0.825	641
## 11	0.875	312
## 12	0.925	156
## 13	0.975	316

Summary of Posterior Predictive Histogram

1. **Predicted Probability Clustering (0.6-0.7):** The highest count of predicted probabilities lies in the 0.6–0.7 range, specifically around midpoints like 0.61, 0.63, and 0.65. This concentration suggests that the model favors assigning moderate-to-high stroke probabilities across the sample. This could mean that, for the majority of individuals, the model leans toward assigning a moderate risk of stroke rather than spreading predictions more evenly across lower and higher ranges.
2. **Sparse Low-Probability Predictions (0.51-0.55):** There are few predictions in the low-probability range, especially around 0.5 or lower. This suggests that the model is relatively less confident in assigning low stroke probabilities, possibly indicating that the predictors push the risk estimates upward on average.
3. **Significant High-Probability Predictions (0.87-0.99):** Although there are fewer cases in the extreme high-probability ranges (like 0.9 and above), these counts still stand out. The model produces a fair number of predictions with very high stroke probabilities, suggesting that it is confident in classifying certain cases as high-risk. This distribution could mean that the model is reacting strongly to certain predictors or combinations of predictors that correlate with high stroke risk.
4. **Comparison to Actual Stroke Rate:** The observed stroke rate in this dataset is approximately 4.87% (actual mean of 0.0487), which is much lower than what the model predicts on average. This discrepancy suggests that the model may overestimate stroke risk, likely leaning toward higher probability predictions across the board. This can happen if certain predictors exert a strong influence, pushing the average prediction upward.
5. **Potential Adjustments:**
 - o **Calibration:** Adjusting model calibration could help bring predicted probabilities closer to the observed rate. This might involve refining the priors on influential predictors to moderate their effect.

- **Diagnostics:** Further diagnostic analysis on predictor influence could help isolate variables that lead to higher-than-expected predictions, enabling a more balanced probability distribution.

Observation

Overall, the model appears to overestimate stroke probability for this dataset, which could affect its reliability. Calibrating the model and re-evaluating predictor influence might help achieve more realistic risk predictions in line with the observed stroke occurrence rate.

.

1.12. CONCLUSION AND FURTHER STUDY

1.12.1 CONCLUSION

This concentration of moderate-to-high risk predictions points to the potential value of the model in clinical decision-making. By providing a risk profile, the model highlights individuals who may benefit from targeted screening and preventive interventions. Such an approach can support healthcare providers in implementing proactive measures to mitigate stroke risks in high-risk populations. Overall, the model demonstrates promise as a predictive tool in healthcare settings, where accurate risk assessments are essential for personalized prevention strategies. Most predictions fall within this range, demonstrating that the model is not only sensitive to the risk factors included but also capable of providing actionable insights regarding stroke risk (Katan & Luft, 2018; Hankey & Smith, 2020).

The observed distribution highlights a critical insight: a substantial proportion of individuals are predicted to have a significant probability of stroke, underscoring the importance of targeted screening and preventive measures in clinical settings. The model's ability to accurately capture the likelihood of stroke occurrence supports its utility in guiding healthcare decisions and interventions aimed at high-risk populations (Smith & Benjamin, 2019).

1.12.2 Recommendations for Further Study

- **Threshold Analysis:** Conduct a detailed analysis to establish optimal thresholds for classifying individuals as high-risk based on predicted probabilities. This could involve examining trade-offs between sensitivity and specificity and determining the implications for clinical interventions.
- **Comparison with Actual Outcomes:** Validate the model's predictions against actual stroke occurrences in the dataset. This will provide insights into the model's predictive accuracy and reliability, allowing for adjustments and refinements if necessary.

- **Inclusion of Additional Predictors:** Explore the potential benefits of including additional predictors or interaction terms in the model. Factors such as family history of stroke, medication use, or lifestyle changes could enhance the predictive power and provide a more comprehensive risk assessment.
- **Longitudinal Study Design:** Consider a longitudinal approach to track individuals over time and observe actual stroke occurrences in relation to predicted probabilities. This will help assess the model's effectiveness in a real-world setting and provide insights into the dynamic nature of stroke risk factors.
- **External Validation:** Test the model on external datasets to evaluate its generalizability and robustness. This is crucial for understanding how well the model performs across different populations and settings.
- **Patient Stratification:** Investigate how predicted probabilities can inform personalized treatment strategies. Understanding the different risk profiles may help in developing tailored interventions for various subgroups of patients.
- **Integration of Machine Learning Techniques:** Explore the application of machine learning methods in conjunction with Bayesian modeling to enhance predictive accuracy. This could include ensemble methods or deep learning approaches that may capture complex relationships within the data. By addressing these areas in future research, the understanding of stroke risk can be deepened, ultimately leading to improved prevention strategies and patient outcomes.

2. REFERENCES

- Feigin, V.L., Norrving, B. and Mensah, G.A., 2014. Global burden of stroke. *Circulation Research*, 114(9), pp. 145-149.
- Hankey, G.J., & Smith, E., 2020. Epidemiology of stroke. *Lancet*, 392(10154), pp. 29-40. DOI: 10.1016/S0140-6736(20)31577-1.
- Katan, M., & Luft, A., 2018. Global burden of stroke. *The Lancet Neurology*, 17(4), pp. 251-252. DOI: 10.1016/S1474-4422(18)30061-4.
- Norrving, B., & Kissela, B., 2016. The challenge of stroke: A global perspective. *Stroke*, 47(4), pp. e73-e80. DOI: 10.1161/STROKEAHA.115.010640.
- Pannala, A.S. & Reddy, V., 2018. Association between fasting blood glucose and the risk of stroke: A systematic review and meta-analysis. *European Journal of Neurology*, 25(9), pp. 1134-1141. DOI: 10.1111/ene.13757.
- Smith, S.C. & Benjamin, E.J., 2019. AHA/ACC guidelines for the prevention of stroke. *Circulation*, 140(9), pp. e778-e796. DOI: 10.1161/CIR.0000000000000466.
- World Health Organization, 2021. Stroke. [online] Available at: <https://www.who.int/news-room/fact-sheets/detail/stroke> [Accessed 8 October 2024].
- <https://www.kaggle.com/datasets/fedesoriano/stroke-prediction-dataset>